

1 Introduction

Compilers translate programs written in a source language (e.g., Java) into semantically equivalent programs in some target language (e.g., assembly code). They let us make our source language arbitrarily abstract so we can write programs in ways that humans understand while letting the computer execute programs in ways that machines understand. In a perfect world, such translation would be straightforward. Reality, however, is unforgiving. Straightforward compilation results in clunky target code that performs a lot of redundant computations. To produce efficient code, we must rely on less-than-straightforward methods. Typical compilers go through a stage of *optimization*, whereby a number of semantics-preserving transformations are applied to an *intermediate representation* of the source code. These then (hopefully) produce a more efficient version of said representation. Optimizers tend to work in *phases*, applying specific transformations during any given phase.

Global value numbering (GVN) is such an analysis performed by many highly-optimizing compilers. Its roots run deep through both the theoretical and the practical. Using the results of this analysis, the compiler can identify expressions in the source code that produce the same value—not just by lexical comparison (i.e., variables having the same name), but by proving equivalences between what’s actually computed at runtime. These expressions can then be simplified by further algorithms for redundancy elimination. This is the very essence of most compiler optimizations: avoid redundant computation, giving us code that runs as quickly as possible while still following what the programmer originally wrote.

High-level, dynamic languages tend to suffer from efficiency issues: they’re often interpreted rather than compiled, and perform no heavy optimization of the source code. However, the Factor language (<http://factorcode.org>) fills an intriguing design niche, as it’s very high-level yet still fully compiled. It’s still young, though, so its compiler craves all the improvements it can get. In particular, while Factor currently has a *local* value numbering analysis, it is inferior to GVN in several significant ways.

In this thesis, we explore the implementation and use of GVN in improving the strength of optimizations in Factor. Because Factor is a young and relatively unknown language, Section 2 provides a short tutorial, laying a foundation for understanding the changes. Section 3 describes the overall architecture of the Factor compiler, highlighting where the exact contributions of this thesis fit in. Finally, Section 4 goes into detail about the existing and new value numbering passes, closing with a look at the results achieved and directions for future work.

All the code for the GVN phase was written atop Factor version 0.94, and a copy of it can be found in the appendix. In the unlikely event that you want to cite this thesis, you may use the following BIBTEX entry:

```
@mastersthesis{vondrak:11,
  author = {Alex Vondrak},
  title  = {Global Value Numbering in Factor},
  school = {California Polytechnic State University, Pomona},
  month  = sep,
  year   = {2011},
}
```

ref

make sure
it's single-
page




Figure 1: Visualizing stack-based calculation

2 Language Primer

citations for this history are fragmented across the internet

Factor is a rather young language created by Slava Pestov in September of 2003. Its first incarnation targeted the Java Virtual Machine (JVM) as an embedded scripting language for a game. As such, its feature set was minimal. Factor has since evolved into a general-purpose programming language, gaining new features and redesigning old ones as necessary for larger programs. Today’s implementation sports an extensive standard library and has moved away from the JVM in favor of native code generation. In this section, we cover the basic syntax and semantics of Factor for those unfamiliar with the language. This should be just enough to understand the later material in this thesis. More thorough documentation can be found via Factor’s website, <http://factorcode.org>.

2.1 Stack-Based Languages

Like Reverse Polish Notation (RPN) calculators, Factor’s evaluation model uses a global stack upon which operands are pushed before operators are called. This naturally facilitates postfix notation, in which operators are written after their operands. For example, instead of $1 + 2$, we write `1 2 +`. Figure 1 shows how `1 2 +` works conceptually:

- 1 is pushed onto the stack
- 2 is pushed onto the stack
- `+` is called, so two values are popped from the stack, added, and the result (`3`) is pushed back onto the stack

Other stack-based programming languages include Forth, Joy, Cat, and PostScript.

The strength of this model is its simplicity. Evaluation essentially goes left-to-right: literals (like `1` and `2`) are pushed onto the stack, and operators (like `+`) perform some computation using values currently on the stack. This “flatness” makes parsing easier, since we don’t need complex grammars with subtle ambiguities and precedence issues. Rather, we basically just scan left-to-right for tokens separated by whitespace. In the Forth tradition, functions are called *words* since they’re made up of any contiguous non-whitespace characters. This also lends to the term *vocabulary* instead of “module” or “library”. In Factor, the parser works as follows.

- If the current character is a double-quote (`"`), try to parse ahead for a string literal.
- Otherwise, scan ahead for a single token.

cite

cite

cite

cite




Figure 2: Quotations

- If the token is the name of a *parsing word*, that word is invoked with the parser’s current state.
- If the token is the name of an ordinary (i.e., non-parsing) word, that word is added to the parse tree.
- Otherwise, try to parse the token as a numeric literal.

Parsing words serve as hooks into the parser, letting Factor users extend the syntax dynamically. For instance, instead of having special knowledge of comments built into the parser, the parsing word `!` scans forward for a newline and discards any characters read (adding nothing to the parse tree).

Similarly, there are parsing words for what might otherwise be hard-coded syntax for data structure literals. Many act as sided delimiters: the parsing word for the left-delimiter will parse ahead until it reaches the right-delimiter, using whatever was read in between to add objects to the data structure. For example, `{ 1 2 3 }` denotes an array of three numbers. Note the deliberate spaces in between the tokens, so that the delimiters are themselves distinct words. In `{_1_2_3_}` (with spaces as marked), the parsing word `{` parses objects until it reaches `}`, collecting the results into an array. The `{` word would not be called if not for that space, whereas `{1_2_3}` parses as the word `{1`, the number `2`, and the word `3}`—not an array. Further, since the left-delimiter words parse recursively, such literals can be nested, contain comments, etc. Other literals include the following.

<code>V{ 1 2 3 }</code>	<code>! vector</code>
<code>B{ 1 2 3 }</code>	<code>! byte array</code>
<code>BV{ 1 2 3 }</code>	<code>! byte vector</code>
<code>HS{ 1 2 3 }</code>	<code>! hash set</code>
<code>H{ { key1 val1 } { key2 val2 } }</code>	<code>! hash table</code>

Listing 1: Data structure literals in Factor

A particularly important set of parsing words in Factor are the square brackets, `[` and `]`. Any code in between such brackets is collected up into a special sequence called a *quotation*. Essentially, it’s a snippet of code whose execution is suppressed. The code inside a quotation can then be run with the `call` word. Quotations are like anonymous functions in other languages, but the stack model makes them conceptually simpler, since we don’t have to worry about variable binding and the like. Consider a small example like `1 2 [+] call`. You can think of `call` working by “erasing” the brackets around a quotation, so this example behaves just like `1 2 +`. Figure 2 shows its evaluation: instead of adding the numbers immediately, `+` is placed in a quotation, which is

pushed to the stack. The quotation is then invoked by **call**, so **+** pops and adds the two numbers and pushes the result onto the stack. We'll show how quotations are used in Section 2.5.

2.2 Stack Effects

Everything else about Factor follows from the stack-based structure outlined in Section 2.1. Consecutive words transform the stack in discrete steps, thereby shaping a result. In a way, words are functions from stacks to stacks—from “before” to “after”—and whitespace is effectively function composition. Even literals (numbers, strings, arrays, quotations, etc.) can be thought of as functions that take in a stack and return that stack with an extra element pushed onto it.

With this in mind, Factor requires that the number of elements on the stack (the *stack height*) is known at each point of the program in order to ensure consistency. To this end, every word is associated with a *stack effect* declaration using a notation implemented by parsing words. In general, a stack effect declaration has the form

```
( input1 input2 ... -- output1 output2 ... )
```

where the parsing word **(** scans forward for the special token **--** to separate the two sides of the declaration, and then for the **)** token to end the declaration. The names of the intermediate tokens don't technically matter—only how many of them there are. However, names should be meaningful for clarity's sake. The number of tokens on the left side of the declaration (before the **--**) indicates the minimum stack height expected before executing the word. Given exactly this number of inputs, the number of tokens on the right side is the stack height after executing the word.

For instance, the stack effect of the **+** word is **(x y -- z)**, as it pops two numbers off the stack and pushes one number (their sum) onto the stack. This could be written any number of ways, though. **(x x -- x)**, **(number1 number2 -- sum)**, and **(m n -- m+n)** are all equally valid. Further, while the stack effect **(junk x y -- junk z)** has the same relative height change, this declaration would be wrong, since **+** might legitimately be called on only two inputs.

For the purposes of documentation, of course, the names in stack effects do matter. They correspond to elements of the stack from bottom-to-top. So, the rightmost value on either side of the declaration names the top element of the stack. We can see this in Figure 3 on the next page, which shows the effects of standard *stack shuffler* words. These words are used for basic data flow in Factor programs. For example, to discard the top element of the stack, we use the **drop** word, whose effect is simply **(x --)**. To discard the element just below the top of the stack, we use **nip**, whose effect is **(x y -- y)**. This stack effect indicates that there are at least two elements on the stack before **nip** is called: the top element is **y**, and the next element is **x**. After calling the word, **x** is removed, leaving the original **y** still on top of the stack. Other shuffler words that remove data from the stack are **2drop** with the effect **(x y --)**, **3drop** with the effect **(x y z --)**, and **2nip** with the effect **(x y z -- z)**.

The next stack shufflers duplicate data. **dup** copies the top element of the stack, as indicated by its effect **(x -- x x)**. **over** has the effect **(x y -- x y x)**, which tells us that it expects at least two inputs: the top of the stack is **y**, and the next object is **x**. **x** is copied and pushed on top of the two original elements, sandwiching **y** between two **xs**. Other shuffler words that

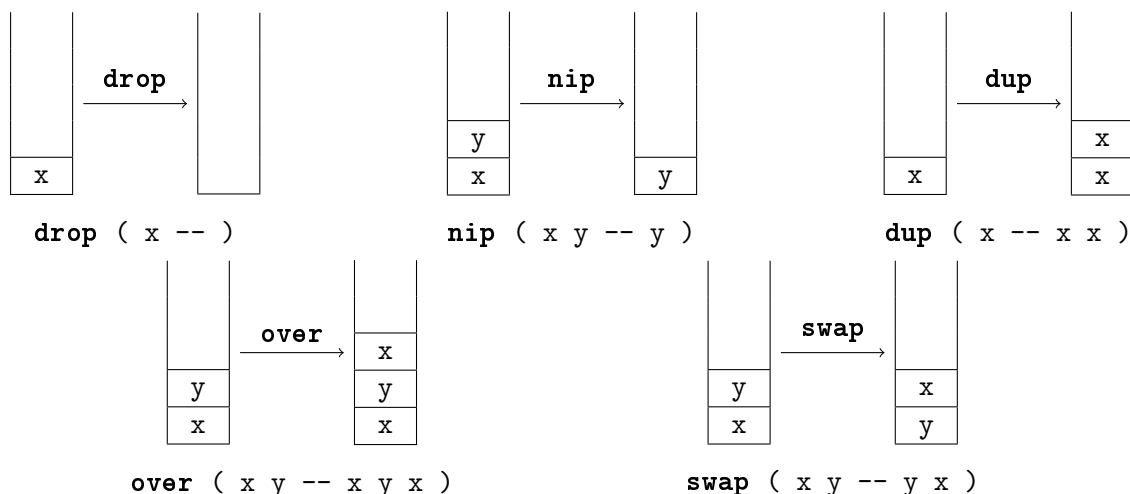


Figure 3: Stack shuffler words and their effects

duplicate data on the stack are **2dup** with the effect $(x\ y\ --\ x\ y\ x\ y)$, **3dup** with the effect $(x\ y\ z\ --\ x\ y\ z\ x\ y\ z)$, **2over** with the effect $(x\ y\ z\ --\ x\ y\ z\ x\ y)$, and **pick** with the effect $(x\ y\ z\ --\ x\ y\ z\ x)$.

True to the name **swap**, the final shuffler in Figure 3 permutes the top two elements of the stack, reversing their order. The stack effect $(x\ y\ --\ y\ x)$ indicates as much. The left side denotes that two inputs are on the stack (the top is *y*, the next is *x*), and the right side shows the outputs are swapped (the top element is *x* and the next is *y*). Factor has other words that permute elements deeper into the stack. However, their use is discouraged because it's harder for the programmer to mentally keep track of more than a couple items on the stack. We'll see how more complex data flow patterns are handled in Section 2.5.2.

2.3 Definitions

```
: hello-world ( -- )
  "Hello, world!" print ;
```

Listing 2: Hello World in Factor

```
: norm ( x y -- norm )
  dup * swap dup * + sqrt ;
```

Listing 3: The Euclidean norm, $\sqrt{x^2 + y^2}$

Using the basic syntax of stack effect declarations described in Section 2.2, we can now understand how to define words. Most words are defined with the parsing word **:**, which scans for a name, a stack effect, and then any words up until the **;** token, which together become the body of the definition. Thus, the classic example in Listing 2 defines a word named **hello-world** which




Figure 4: `norm` example

expects no inputs and pushes no outputs onto the stack. When called, this word will display the canonical greeting on standard output using the `print` word.

A slightly more interesting example is the `norm` word in Listing 3 on the preceding page. This squares each of the top two numbers on the stack, adds them, then takes the square root of the sum. Figure 4 shows this in action. By defining a word to perform these steps, we can replace virtually any instance of `dup * swap dup * + sqrt` in a program simply with `norm`. This is a deceptively important point. Data flow is made explicit via stack manipulation rather than being hidden in variable assignments, so repetitive patterns become painfully evident. This makes identifying, extracting, and replacing redundant code easy. Often, you can just copy a repetitive sequence of words into its own definition verbatim. This emphasis on “factoring” your code is what gives Factor its name.

```

: ^2 ( n -- n^2 )
  dup * ;

: norm ( x y -- norm )
  ^2 swap ^2 + sqrt ;

```

Listing 4: `norm` refactored

As a simple case in point, we see the subexpression `dup *` appears twice in the definition of `norm` in Listing 3 on the preceding page. We can easily factor that out into a new word and substitute it for the old expressions, as in Listing 4. By contrast, programs in more traditional languages are laden with variables and syntactic noise that require more work to refactor: identifying free variables, pulling out the right functions without causing finicky syntax errors, calling a new function with the right variables, etc. Though Factor’s stack-based paradigm is atypical, it is part of a design philosophy that aims to facilitate readable code focusing on short, reusable definitions.

2.4 Object Orientation

You may have noticed that the examples in Section 2.3 did not use type declarations. While Factor is dynamically typed for the sake of simplicity, it does not do away with types altogether. In fact, Factor is object-oriented. However, its object system doesn't rely on classes possessing particular methods, as is common. Instead, it uses *generic words* with methods implemented for particular classes. To start, though, we must see how classes are defined.

2.4.1 Tuples

```
TUPLE: class
    slot-spec1 slot-spec2 slot-spec3 ... ;

TUPLE: subclass < superclass
    slot-spec1 slot-spec2 slot-spec3 ... ;
```

Listing 5: Basic tuple definition syntax

The central data type of Factor's object system is called a *tuple*, which is a class composed of named *slots*—like instance variables in other languages. Tuples are defined with the **TUPLE:** parsing word as shown in Listing 5. A class name is specified first; if it is followed by the **<** token and a superclass name, the tuple inherits the slots of the superclass. If no superclass is specified, the default is the **tuple** class.

Slots can be specified in several ways. The simplest is to just provide a single token, which is the name of the slot. This slot can then hold any type of object. Using the syntax { **name class** }, a slot can be limited to hold only instances of a particular class, like **integer** or **string**. There are other forms of slot specifiers, which we will cover after some examples.

```
TUPLE: color ;

: <color> ( -- color )
    color new ;

TUPLE: rgb < color red green blue ;

: <rgb> ( r g b -- rgb )
    rgb boa ;
```

Listing 6: Simple tuple examples

Consider the two tuples defined in Listing 6. The first, **color**, has no slots. With every tuple, a class predicate is defined with the stack effect (**object -- ?**) whose name is the class suffixed by a question mark. Here, the word **color?** is defined, which pushes a boolean (in Factor, either **t** or **f**) indicating whether the top element of the stack is an instance of the **color** class. The second tuple, **rgb**, inherits from the **color** class. While this doesn't give **rgb** any different slots, it

does mean that an instance of `rgb` will return `t` for `color`? due to the “is-a” relationship between subclass and superclass. The word `rgb?` is similarly defined.

Notice that the `rgb` tuple declares three slots named `red`, `green`, and `blue`. Since the slots’ classes aren’t declared, any sort of object can be stored in them. A set of methods are defined to manipulate an `rgb` instance’s slots. Three *reader* words are defined (one for each slot), analogous to “getter” methods in other languages. Following the template for naming reader words, this example defines `red>>`, `green>>`, and `blue>>`. Each word has the stack effect (`object -- value`), and extracts the value corresponding to the eponymous slot. Similarly, the *writer* words `red<<`, `green<<`, and `blue<<` each have the stack effect (`value object --`), and store values in the corresponding `rgb` slots destructively. To leave the modified `rgb` instance on the stack while setting slots, the *setter* words `>>red`, `>>green`, and `>>blue` are also defined, each with the stack effect (`object value -- object'`). These words are defined in terms of writers. For instance, `>>red` is the same as `over red<<`, since `over` copies a reference to the tuple (i.e., it doesn’t make a “deep” copy).

To construct an instance of a tuple, we can use either `new` or `boa`. `new` will not initialize any of the slots to a particular input value—all slots will default to Factor’s canonical false value, `f`. `new` is used in Listing 6 on the preceding page to define `<color>` (by convention, the constructor for `foo` is named `<foo>`). First, we push the class `color` onto the stack (this word is also automatically defined by `TUPLE:`), then just call `new`, leaving a new instance on the stack. Since this particular tuple has no slots, using `new` makes sense. We might also use it to initialize a class, then use setter words to only assign a particular subset of slots’ values.

However, we often want to initialize a tuple with values for each of its slots. For this, we have `boa`, which works similarly to `new`. This is used in the definition of `<rgb>` in Listing 6 on the previous page. The difference here is the additional inputs on the stack—one for each slot, in the order they’re declared. That is, we’re constructing the tuple **by order of arguments**, giving us the fun pun “`boa` constructor”. So, `1 2 3 <rgb>` will construct an `rgb` instance with the `red` slot set to 1, the `green` slot set to 2, and the `blue` slot set to 3.

```
TUPLE: email
{ from string }
{ to array }
{ cc array }
{ bcc array }
{ subject string }
{ content-type string initial: "text/plain" }
{ encoding word initial: utf8 }
{ body string } ;
```

Listing 7: Special slot specifiers

Now that we’ve seen the various words defined for tuples, we can explore more complex slot specifiers. Using the array-like syntax from before, slot specifiers may be marked with certain *attributes*—both with the class declared (like { `name class attributes...` }) and without the class declared (as in { `name attributes...` }). In particular, Factor recognizes two different attributes. If a slot marked **read-only**, the writer (and thus setter) for the slot will not be defined,

so the slot cannot be altered. A slot may also provide an initial value using the syntax `initial:some-literal`. This will be the slot’s value when instantiated with `new`.

For example, Listing 7 on the preceding page shows a tuple definition from Factor’s `smtp` vocabulary that defines an `email` object. The `from` address, `subject`, and `body` must be instances of `string`, while `to`, `cc`, and `bcc` are `arrays` of destination addresses. The `content-type` slot must also be a `string`, but if unspecified, it defaults to `"text/plain"`. The `encoding` must be a `word` (in Factor, even words are first-class objects), which by default is `utf8`, a word from the `io.encodings.utf8` vocabulary for a Unicode format.

2.4.2 Generics and Methods

Unlike more common object systems, we do not define individual methods that “belong” to particular tuples. In Factor, you define a method that specializes on a class for a particular generic word. That way, when the generic word is called, it dispatches on the class of the object, invoking the most specific method for the object.

Generic words are declared with the syntax `GENERIC: word-name (stack -- effect)`. Words defined this way will then dispatch on the class of the top element of the stack (necessarily the rightmost input in the stack effect). To define a new method with which to control this dispatch, we use the syntax `M: class word-name definition... ;`.

```
USING: bit-sets hash-sets sequences ;
IN: sets

MIXIN: set
INSTANCE: sequence set
INSTANCE: hash-set set
INSTANCE: bit-set set
```

Listing 8: Set instances

An accessible example of a generic word is in Factor’s `sets` vocabulary. `set` is a *mixin* class—a union of other classes whose members may be extended by the user. We can see the standard definition in Listing 8. Note that the `USING:` form specifies vocabularies being used (like Java’s `import`), and `IN:` specifies the vocabulary in which the definitions appear (like Java’s `package`). We can see here that instances of the `sequence`, `hash-set`, and `bit-set` classes are all instances of `set`, so will respond `t` to the predicate `set?`. Similarly, `sequence` is a mixin class with many more members, including `array`, `vector`, and `string`.

Listing 9 on the following page shows the `cardinality` generic from Factor’s `sets` vocabulary, along with its methods. This generic word takes a `set` instance from the top of the stack and pushes the number of elements it contains. For instance, if the top element is a `bit-set`, we extract its `table` slot and invoke another word, `bit-count`, on that. But if the top element is `f` (the canonical false/empty value), we know the cardinality is 0. For any `sequence`, we may offshore the work to a different generic, `length`, defined in the `sequences` vocabulary. The final method gives a default behavior for any other `set` instance, which simply uses `members` to obtain an equivalent `sequence` of set members, then calls `length`.

```

IN: sets
GENERIC: cardinality ( set -- n )

USING: accessors bit-sets math.bitwise sets ;
M: bit-set cardinality table>> bit-count ;

USING: kernel sets ;
M: f cardinality drop 0 ;

USING: accessors assocs hash-sets sets ;
M: hash-set cardinality table>> assoc-size ;

USING: sequences sets ;
M: sequence cardinality length ;

USING: sequences sets ;
M: set cardinality members length ;

```

Listing 9: Set cardinality using Factor’s object system

By viewing a class as a set of all objects that respond positively to the class predicate, we may partially order classes with the subset relationship. Method dispatch will use this ordering when **cardinality** is called to select the most specific method for the object being dispatched upon. For instance, while no explicit method for **array** is defined, any instance of **array** is also an instance of **sequence**. In turn, every instance of **sequence** is also an instance of **set**. We have methods that dispatch on both **set** and **sequence**, but the latter is more specific, so that is the method invoked. If we define our own class, **foo**, and declare it as an instance of **set** but not as an instance of **sequence**, then the **set** method of **cardinality** will be invoked. Sometimes resolving the precedence gets more complicated, but these edge-cases are beyond the scope of our discussion.

2.5 Combinators

Quotations, introduced in Section 2.1, form the basis of both control flow and data flow in Factor. Not only are they the equivalent of anonymous functions, but the stack model also makes them syntactically lightweight enough to serve as blocks akin to the code between curly braces in C or Java. Higher-order words that make use of quotations on the stack are called *combinators*. It’s simple to express familiar conditional logic and loops using combinators, as we’ll show in Section 2.5.1. In the presence of explicit data flow via stack operations, even more patterns arise that can be abstracted away. Section 2.5.2 explores how we can use combinators to express otherwise convoluted stack-shuffling logic more succinctly.

2.5.1 Control Flow

The most primitive form of control flow in typical programming languages is, of course, the **if** statement, and the same holds true for Factor. The only difference is that Factor’s **if** isn’t

```

5 even? [ "even" print ] [ "odd" print ] if

{ } empty? [ "empty" print ] [ "full" print ] if

100 [ "isn't f" print ] [ "is f" print ] if

```

Listing 10: Conditional evaluation in Factor

syntactically significant—it’s just another word, albeit implemented as a primitive. For the moment, it will do to think of **if** as having the stack effect (`? true false --`). The third element from the top of the stack is a condition, and it’s followed by two quotations. The first quotation (second element from the top of the stack) is called if the condition is true, and the second quotation (the top of the stack) is called if the condition is false. Specifically, **f** is a special object in Factor for falsity. It is a singleton object—the sole instance of the **f** class—and is the only false value in the entire language. Any other object is necessarily boolean true. For a canonical boolean, there is the **t** object, but its truth value exists only because it is not **f**. Basic **if** use is shown in Listing 10. The first example will print “odd”, the second “empty”, and the third “isn’t f”. All of them leave nothing on the stack.

vref

```

: example1 ( x -- 0/x-1 )
  dup even? [ drop 0 ] [ 1 - ] if ;

: example2 ( x y -- x+y/x-y )
  2dup mod 0 = [ + ] [ - ] if ;

: example3 ( x y -- x+y/x )
  dup odd? [ + ] [ drop ] if ;

```

Listing 11: **if**’s stack effect varies

However, the simplified stack effect for **if** is quite restrictive. (`? true false --`) intuitively means that both the **true** and **false** quotations can’t take any inputs or produce any outputs—that their effects are (`--`). We’d like to loosen this restriction, but per Section 2.2, Factor must know the stack height after the **if** call. We could give **if** the effect (`x ? true false -- y`), so that the two quotations could each have the stack effect (`x -- y`). This would work for the **example1** word in Listing 11, yet it’s just as restrictive. For instance, the **example2** word would need **if** to have the effect (`x y ? true false -- z`), since each branch has the effect (`x y -- z`). Furthermore, the quotations might even have different effects, but still leave the overall stack height balanced. Only one item is left on the stack after a call to **example3** regardless, even though the two quotations have different stack effects: **+** has the effect (`x y -- z`), while **drop** has the effect (`x --`).

In reality, there are infinitely many correct stack effects for **if**. Factor has a special notation for such *row-polymorphic* stack effects. If a token in a stack effect begins with two dots, like `..a` or `..b`, it is a *row variable*. If either side of a stack effect begins with a row variable, it represents any number inputs/outputs. Thus, we could give **if** the stack effect

```
( ..a ? true false -- ..b )
```

to indicate that there may be any number of inputs below the condition on the stack, and any number of outputs will be present after the call to **if**. Note that these numbers aren't necessarily equal, which is why we use distinct row variables in this case. However, this still isn't quite enough to capture the stack height requirements. It doesn't communicate that **true** and **false** must affect the stack in the same ways. For this, we can use the notation **quot: (stack -- effect)**, giving quotations a nested stack effect. Using the same names for row variables in both the “inner” and “outer” stack effects will refer to the same number of inputs or outputs. Thus, our final (correct) stack effect for **if** is

```
( ..a ? true: ( ..a -- ..b ) false: ( ..a -- ..b ) -- ..b )
```

This tells us that the **true** quotation and the **false** quotation will each create the same relative change in stack height as **if** does overall.

```
{ "Lorem" "ipsum" "dolor" } [ print ] each

0 { 1 2 3 } [ + ] each

10 iota [ number>string print ] each

3 [ "Ho!" print ] times

[ t ] [ "Infinite loop!" print ] while

[ f ] [ "Executed once!" print ] do while
```

Listing 12: Loops in Factor

Though **if** is necessarily a language primitive, other control flow constructs are defined in Factor itself. It's simple to write combinators for iteration and looping as tail-recursive words that invoke quotations. Listing 12 showcases some common looping patterns. The most basic yet versatile word is **each**. Its stack effect is

```
( ... seq quot: ( ... x -- ... ) -- ... )
```

Each element **x** of the sequence **seq** will be passed to **quot**, which may use any of the underlying stack elements. Here, unlike **if**, we enforce that the input stack height is exactly the same as the output (since we use the same row variable). Otherwise, depending on the number of elements in **seq**, we might dig arbitrarily deep into the stack or flood it with a varying number of values. The first use of **each** in Listing 12 is balanced, as the quotation has the effect **(str --)** and no additional items were on the stack to begin with. Essentially, it's equivalent to **"Lorem" print "ipsum" print "dolor" print**. On the other hand, the quotation in the second example has the stack effect **(total n -- total+n)**. This is still balanced, since there is one additional item

below the sequence on the stack (namely 0), and one element is left by the end (the sum of the sequence elements). So, this example is the same as `0 1 + 2 + 3 +`.

Any instance of the extensive **sequence** mixin will work with **each**, making it very flexible. The third example in Listing 12 on the preceding page shows **iota**, which is used here to create a *virtual* sequence of integers from 0 to 9 (inclusive). No actual sequence is allocated, merely an object that behaves like a sequence. In Factor, it's common practice to use **iota** and **each** in favor of repetitive C-like **for** loops.

Of course, we sometimes don't need the induction variable in loops. That is, we just want to execute a body of code a certain number of times. For these cases, there's the **times** combinator, with the stack effect

```
( ... n quot: ( ... -- ... ) -- ... )
```

This is similar to **each**, except that **n** is a number (so we needn't use **iota**) and the quotation doesn't expect an extra argument (i.e., a sequence element). Therefore, the example in Listing 12 on the previous page is equivalent to `"Ho!" print "Ho!" print "Ho!" print`.

Naturally, Factor also has the **while** combinator, whose stack effect is

```
( ..a pred: ( ..a -- ..b ? ) body: ( ..b -- ..a ) -- ..b )
```

The row variables are a bit messy, but it works as you'd expected: the **pred** quotation is invoked on each iteration to determine whether **body** should be called. The **do** word is a handy modifier for **while** that simply executes the body of the loop once before leaving **while** to test the precondition as per usual. Thus, the last example in Listing 12 on the preceding page executes the body once, despite the condition being immediately false.

```
{ 1 2 3 } [ 1 + ] map
{ 1 2 3 4 5 } [ even? ] filter
{ 1 2 3 } 0 [ + ] reduce
```

Listing 13: Higher-order functions in Factor

In the preceding combinators, quotations were used like blocks of code. But really, they're the same as anonymous functions from other languages. As such, Factor borrows classic tools from functional languages, like **map** and **filter**, as shown in Listing 13. **map** is like **each**, except that the quotation should produce a single output. Each such output is collected up into a new sequence of the same class as the input sequence. Here, the example produces `{ 2 3 4 }`. **filter** selects only those elements from the sequence for which the quotation returns a true value. Thus, the **filter** in Listing 13 outputs `{ 2 4 }`. Even **reduce** is in Factor, also known as a *left fold*. An initial element is iteratively updated by pushing a value from the sequence and invoking the quotation. In fact, **reduce** is defined as **swapped each**, where **swapped** is a shuffler word with the stack effect `(x y z -- y x z)`. Thus, the example in Listing 13 is the same as `0 { 1 2 3 } [+] each`, as in Listing 12 on the preceding page.

These are just some of the control flow combinators defined in Factor. Several variants exist that meld stack shuffling with control flow, or can be used to shorten common patterns like empty false branches. An entire list is beyond the scope of our discussion, but the ones we've studied should give a solid view of what standard conditional execution, iteration, and looping looks like in a stack-based language.

2.5.2 Data Flow

While avoiding variables and additional syntax makes it easier to refactor code, keeping mental track of the stack can be taxing. If we need to manipulate more than the top few elements of the stack, code gets harder to read and write. Since the flow of data is made explicit via stack shufflers, we actually wind up with redundant patterns of data flow that we otherwise couldn't identify. In Factor, there are several combinators that clean up common stack-shuffling logic, making code easier to understand.

```
: without-dip1 ( x y -- x+1 y )
  swap 1 + swap ;

: with-dip1 ( x y -- x+1 y )
  [ 1 + ] dip ;

: without-dip2 ( x y z -- x-y z )
  2over - swapd nip swapd nip swap ;

: with-dip2 ( x y z -- x-y z )
  [ - ] dip ;

: without-keep1 ( x -- x+1 x )
  dup 1 + swap ;

: with-keep1 ( x -- x+1 x )
  [ 1 + ] keep ;

: without-keep2 ( x y -- x-y y )
  swap over - swap ;

: with-keep2 ( x y -- x-y y )
  [ - ] keep ;
```

Listing 14: Preserving combinators

The first combinators we'll look at are **dip** and **keep**. These are used to preserve elements of the stack. When working with several values, sometimes we don't want to use all of them at quite the same time. Using **drop** and the like wouldn't help, as we'd lose the data altogether. Rather, we want to retain certain stack elements, do a computation, then restore them. For an unconvincing

but illustrative example, suppose we have two values on the stack, but we want to increment the second element from the top. **without-dip1** in Listing 14 on the previous page shows one strategy, where we shuffle the top element away with **swap**, perform the computation, then **swap** the top back to its original place. A cleaner way is to call **dip** on a quotation, which will execute that quotation just under the top of the stack, as in **with-dip1**. While the stack shuffling in **without-dip1** isn't terribly complicated, it doesn't convey our meaning very well. Shuffling the top element out of the way becomes increasingly difficult with more complex computations. In **without-dip2**, we want to call **-** on the two elements below the top. For lack of a more robust stack shuffler, we use **2over** to isolate the two values so we can call **-**. The rest of the word consists of shuffling to get rid of excess values on the stack. It's also worth noting that **swapped** is a deprecated word in Factor, since its use starts making code harder to reason about. Alternatively, we could dream up a more complex stack shuffler with exactly the stack effect we wanted in this situation. But this solution doesn't scale: what if we had to calculate something that required more inputs or produced more outputs? Clearly, **dip** provides a cleaner alternative in **with-dip2**.

keep provides a way to hold onto the top element of the stack, but still use it to perform a computation. In general, `[...] keep` is equivalent to `dup [...] dip`. Thus, the current top of the stack remains on top after the use of **keep**, but the quotation is still invoked with that value. In **with-keep1** in Listing 14 on the preceding page, we want to increment the top, but stash the result below. Again, this logic isn't terribly complicated, though **with-keep1** does away with the shuffling. **without-keep2** shows a messier example where a simple **dup** will not save us, as we're using more than just the top element in the call to **-**. Rather, three of the four words in the definition are dedicated to rearranging the stack in just the right way, obscuring the call to **-** that we really want to focus on. On the other hand, **with-keep2** places the subtraction word front-and-center in its own quotation, while **keep** does the work of retaining the top of the stack.

```
TUPLE: coord x y ;

: without-bi ( coord -- norm )
  [ x>> sq ] keep y>> sq + sqrt ;

: with-bi ( coord -- norm )
  [ x>> sq ] [ y>> sq ] bi + sqrt ;

: without-tri ( x -- x+1 x+2 x+3 )
  [ 1 + ] keep [ 2 + ] keep 3 + ;

: with-tri ( x -- x+1 x+2 x+3 )
  [ 1 + ] [ 2 + ] [ 3 + ] tri ;
```

Listing 15: Cleave combinators

The next set of combinators apply multiple quotations to a single value. The most general form of these so-called *cleave* combinators is the word **cleave**, which takes an array of quotations as input, and calls each one in turn on the top element of the stack. Of course, for only a couple of quotations, wrapping them in an array literal becomes cumbersome. The word **bi** exists for the two-quotation case, and **tri** for the three quotations. Cleave combinators are often used to extract

multiple slots from a tuple. Listing 15 on the previous page shows such a case in the `with-bi` word, which improves upon using just `keep` in the `without-bi` word. In general, a series of `keeps` like `[a] keep [b] keep c` is the same as `{ [a] [b] [c] } cleave`, which is more readable. We can see this in action in the difference between `without-tri` and `with-tri` in Listing 15 on the preceding page. In cases where we need to apply multiple quotations to a set of values instead of just a single one, there are also the variants `2cleave` and `3cleave` (and the corresponding `2bi`, `2tri`, `3bi`, and `3tri`), which apply the quotations to the top two and three elements of the stack, respectively.

```
: without-bi* ( str1 str2 -- str1' str2' )
  [ >upper ] dip >lower ;

: with-bi* ( str1 str2 -- str1' str2' )
  [ >upper ] [ >lower ] bi* ;

: without-tri* ( x y z -- x+1 y+2 z+3 )
  [ [ 1 + ] dip 2 + ] dip 3 + ;

: with-tri* ( x y z -- x+1 y+2 z+3 )
  [ 1 + ] [ 2 + ] [ 3 + ] tri* ;
```

Listing 16: Spread combinators

To apply multiple quotations to multiple values, Factor has *spread* combinators. Whereas cleave combinators abstract away repeated instances of `keep`, spread combinators replace nested calls to `dip`. The archetypical combinator, `spread`, takes an array of quotations, like `cleave`. However, instead of applying each one to the top element of the stack, each one corresponds to a separate element of the stack. Thus, `{ [a] [b] } spread` invokes `b` on the top element, and `a` on the element beneath the top. Much like `cleave`, there are shorthand words for the two- and three-quotation cases. These are suffixed with asterisks to indicate the spread variants, so we have `bi*` and `tri*`. In Listing 16, the `without-bi*` word shows the simple `dip` pattern that `bi*` encapsulates. Here, we’re converting the string `str1` (the second element from the top) into uppercase and `str2` (the top element) to lowercase. In `with-bi*`, the `>upper` and `>lower` words are seen first, uninterrupted by an extra word, making the code easier to read. More compelling is the way that `tri*` replaces the `dips` that can be seen in `without-tri*`. In comparison, `with-tri*` is less nested and easier to comprehend at first glance. While there are `2bi*` and `2tri*` variants that spread quotations to two values apiece on the stack, they are uncommon in practice.

Finally, *apply* combinators invoke a single quotation on multiple stack entries in turn. While there is a generalized word, it’s more common to use the corresponding shorthands. Here, they are suffixed with at-signs, so `bi@` applies a quotation to each of the top two stack values, and `tri@` to each of the top three. This way, rather than duplicate code for each time we want to call a word, we need only specify it once. This is demonstrated clearly in Listing 17 on the following page. In `without-bi@`, we see that the quotation `[sq]` (for squaring numbers) appears twice for the call to `bi*`. In general, we can replace spread combinators whose quotations are all the same with a single quotation and an apply combinator. Thus, `with-bi@` cuts down on the duplicated `[sq]` in `without-bi@`. Similarly, we can replace the three repeated quotations passed


```

: without-bi@ ( x y -- norm )
  [ sq ] [ sq ] bi* + sqrt ;

: with-bi@ ( x y -- norm )
  [ sq ] bi@ + sqrt ;

: without-tri@ ( x y z -- x+1 y+1 z+1 )
  [ 1 + ] [ 1 + ] [ 1 + ] tri* ;

: with-tri* ( x y z -- x+1 y+1 z+1 )
  [ 1 + ] tri@ ;

```

Listing 17: Apply combinators

to **tri*** in **without-tri@** with a single instance passed to **tri@** in **with-tri@**. Like other data flow combinators, we have the numbered variants. **2bi@** has the stack effect (**w x y z quot --**), where **quot** expects two inputs, and is thus applied to **w** and **x** first, then to **y** and **z**. Similarly, **2tri@** applies the quotation to the top six objects of the stack in groups of two. Like their spread counterparts, they are not used very much.

Some wrap-up that isn't completely lame.

3 The Factor Compiler

If we could sort programming languages by the fuzzy notions we tend to have about how “high-level” they are, toward the high end we’d find dynamically-typed languages like Python, Ruby, and PHP—all of which are generally more interpreted than compiled. Despite being as high-level as these popular languages, Factor’s implementation is driven by performance. Factor source is always compiled to native machine code using either its simple, non-optimizing compiler or (more typically) the optimizing compiler that performs several sorts of data and control flow analyses. In this section, we look at the general architecture of Factor’s implementation, after which we place a particular emphasis on the transformations performed by the optimizing compiler.

Though there are projects for this

3.1 Organization

At the lowest level, Factor is written atop a C++ virtual machine (VM) that is responsible for basic runtime services. This is where the non-optimizing base compiler is implemented. It’s the base compiler’s job to compile the simplest primitives: operations that push literals onto the data stack, **call**, **if**, **dip**, words that access tuple slots as laid out in memory, stack shufflers, math operators, functions to allocate/deallocate call stack frames, etc. The aim of the base compiler is to generate native machine code as fast as possible. To this end, these primitives correspond to their own stubs of assembly code. Different stubs are generated by Factor depending on the instruction set supported by the particular machine in use. Thus, the base compiler need only make a single pass over the source code, emitting these assembly instructions as it goes.

This compiled code is saved in an *image file*, which contains a complete snapshot of the current state of the Factor instance, similar to many Smalltalk and Lisp systems. As code is parsed and compiled, the image is updated, serving as a cache for compiled code. This modified image can be saved so that future Factor instances needn't recompile vocabularies that are already contained in the image.

cite?

The VM also handles method dispatch and memory management. Method dispatch incorporates a *polymorphic inline cache* to speed up generic words. Each generic word's call site is associated with a state:

- In the *cold* state, the call site's instruction computes the right method for the class being dispatched upon, which is the operation we're trying to avoid. As it does this, a polymorphic inline cache stub is generated, thus transitioning it to the next state.
- In the *inline cache* state, a stub has been generated that caches the locations of methods for classes that have already been seen. This way, if a generic word at a particular call site is invoked often upon only a small number of classes (as is often in the case in loops, for example), we don't need to waste as much time doing method lookup. By default, if more than three different classes are dispatched upon, we transition to the next state.
- In the *megamorphic* state, the call instruction points to a larger cache that is allocated for the specific generic word (i.e., it is shared by all call sites). While not as efficient as an inline cache, this can still improve the performance of method dispatch.

To manage memory, the Factor VM uses a generational garbage collector (GC), which carves out sections of space on the heap for objects of different ages. Garbage in the oldest generation is collected with a mark-sweep-compact algorithm, while younger generations rely on a copying collector. This way, the GC is specialized for large numbers of short-lived objects that will stay in the younger generations without being promoted to the older generation. In the oldest space, even compiled code can be compacted. This is to avoid heap fragmentation in applications that must call the compiler at runtime, such as Factor's interactive development environment.

cite?

Values are referenced by tagged pointers, which use the three least significant bits of the pointer's address to store type information. This is possible because Factor aligns objects on an eight-byte boundary, so the three least significant bits of an address are always 0. These bits give us eight unique tags, but since Factor has more than eight data types, two tags are reserved to indicate that the type information is stored elsewhere. One is for VM types without their own tag, and the other is for user-defined tuples, each of which has its own type. Sufficiently small integers (e.g., 29-bit integers on a 32-bit machine, since the other 3 bits are used for the type tag) are stored directly in the pointer, so they needn't be heap-allocated. Larger integers and floating point numbers are boxed, but the optimizing compiler may unbox them to store floats in registers.

The VM is meant to be minimal, as Factor is mostly *self-hosting*. That is, the real workhorses of the language are written in Factor itself, including the standard libraries, parser, object system, and the optimizing compiler. It's possible for the compiler to be written in Factor because of the *bootstrapping* process that creates a new image from scratch. First, a minimal *boot image* is created from an existing *host* Factor instance. When the VM runs the boot image, it initiates the bootstrapping process. Using the host's parser, the base compiler will compile the core vocabularies

necessary to load the optimizing compiler. Once the optimizing compiler can itself be compiled, it is used to recompile (and thus optimize) all of the words defined so far. With the basic vocabularies recompiled, any additional vocabularies can be loaded using the optimized compiler and saved into a new, working image.

Thus, while the Factor VM is important, it is a small part of the code base. Since the bootstrapping process allows the optimizing compiler (hereafter just “the compiler”) to be written in the same high-level language it’s compiling, we can avoid the fiddly low-level details of the C++ backend. This is more conducive to writing advanced compiler optimizations, which are often complicated enough without having a concise, dynamically-typed, garbage-collected language like Factor to help us.

3.2 High-level Optimizations

To manipulate source code abstractly, we must have at least one intermediate representation (IR)—a data structure representing the instructions. It’s common to convert between several IRs during compilation, as each form offers different properties that facilitate particular analyses. The Factor compiler optimizes code in passes across two different IRs: first at a high-level using the `compiler.tree` vocabulary, then at a low-level with the `compiler.cfg` vocabulary.

The high-level IR arranges code into a vector of `node` objects, which may themselves have children consisting of vectors of `node`—a tree structure that lends to the name `compiler.tree`. This ordered sequence of nodes represents control flow in a way that’s effectively simple, annotated stack code. Listing 18 on the next page shows the definitions of the tuples that represent the “instruction set” of this stack code. Each object inherits (directly or indirectly) from the `node` class, which itself inherits from `identity-tuple`. This is a tuple whose `equal?` method is defined to always return `f` so that no two instances are equivalent unless they are the same instance.

Notice that most nodes define some sort of `in-d` and `out-d` slots, which mark each of them with the input and output data stacks. This represents the flow of data through the program. Here, stack values are denoted simply by integers, giving each value a unique identifier. An `#introduce` instance is inserted wherever the next node requires stack values that have not yet been named. Thus, while `#introduce` has no `in-d`, its `out-d` introduces the necessary stack values. Similarly, `#return` is inserted at the end of the sequence to indicate the final state of the data stack with its `in-d` slot.

The most basic operations of a stack language are, of course, pushing literals and calling functions. The `#push` node thus has a `literal` slot and an `out-d` slot, giving a name to the single element it pushes to the data stack. `#call`, of course, is used for normal word invocations. The `in-d` and `out-d` slots effectively serve as the stack effect declaration. In later analyses, data about the word’s definition may be stored across the `body`, `method`, `class`, and `info` slots.

The word `build-tree` takes a Factor quotation and constructs the equivalent high-level IR form. In Listing 19 on page 21, we see the output of the simple example `[1 +] build-tree`. Note that `T{ class { slot1 value1 } { slot2 value2 } ... }` is the syntax for tuple literals. The first node is a `#push` for the 1 literal. Since `+` needs two input values, an `#introduce` pushes a new “phantom” value. `+` gets turned into a `#call` instance. Notice the `in-d` slot refers to the values in the order that they’re passed to the word, not necessarily the order they’ve been introduced in the

```

TUPLE: node < identity-tuple ;

TUPLE: #introduce < node out-d ;
TUPLE: #return < node in-d info ;

TUPLE: #push < node literal out-d ;
TUPLE: #call < node word in-d out-d body method class info ;

TUPLE: #renaming < node ;
TUPLE: #copy < #renaming in-d out-d ;
TUPLE: #shuffle < #renaming mapping in-d out-d in-r out-r ;

TUPLE: #declare < node declaration ;

TUPLE: #terminate < node in-d in-r ;

TUPLE: #branch < node in-d children live-branches ;
TUPLE: #if < #branch ;
TUPLE: #dispatch < #branch ;

TUPLE: #phi < node phi-in-d phi-info-d out-d terminated ;

TUPLE: #recursive < node in-d word label loop? child ;
TUPLE: #enter-recursive < node in-d out-d label info ;
TUPLE: #call-recursive < node label in-d out-d info ;
TUPLE: #return-recursive < #renaming in-d out-d label info ;

TUPLE: #alien-node < node params ;
TUPLE: #alien-invoke < #alien-node in-d out-d ;
TUPLE: #alien-indirect < #alien-node in-d out-d ;
TUPLE: #alien-assembly < #alien-node in-d out-d ;
TUPLE: #alien-callback < node params child ;

```

Listing 18: High-level IR nodes

```

V{
  T{ #push { literal 1 } { out-d { 6256273 } } }
  T{ #introduce { out-d { 6256274 } } }
  T{ #call
    { word + }
    { in-d V{ 6256274 6256273 } }
    { out-d { 6256275 } }
  }
  T{ #return { in-d V{ 6256275 } } }
}

```

Listing 19: [1 +] build-tree

IR. The sum is pushed to the data stack, so the `out-d` slot is a singleton that names this value. Finally, `#return` indicates the end of the routine, its `in-d` containing the value left on the stack (the sum pushed by `#call`).

```

V{
  T{ #introduce { out-d { 6256132 6256133 } } }
  T{ #shuffle
    { mapping { { 6256134 6256133 } { 6256135 6256132 } } }
    { in-d V{ 6256132 6256133 } }
    { out-d V{ 6256134 6256135 } }
  }
  T{ #return { in-d V{ 6256134 6256135 } } }
}

```

Listing 20: [swap] build-tree

The next tuples in Listing 18 on the previous page reassign existing values on the stack to fresh identifiers. The `#renaming` superclass has the two subclasses `#copy` and `#shuffle`. The former represents the bijection from elements of `in-d` to elements of `out-d` in the same position; corresponding values are copies of each other. The latter represents a more general mapping. Stack shufflers are translated to `#shuffle` nodes with `mapping` slots that dictate how the fresh values in `out-d` correspond to the input values in `in-d`. For instance, Listing 20 shows how **swap** takes in the values 6256132 and 6256133 and outputs 6256134 and 6256135, where the former is mapped to the second element (6256133) and the latter to the first (6256132). Thus, `out-d` swaps the two elements of `in-d`, mapping them to fresh identifiers. The `in-r` and `out-r` slots of `#shuffle` correspond to the *retain* stack, which is an implementation detail beyond the scope of this discussion.

`#declare` is a miscellaneous node used for the `declare` primitive. It simply annotates type information to stack values, as in Listing 21 on the next page. `#terminate` is another one-off node, but a much more interesting one. While Factor normally requires a balanced stack, sometimes we

```
V{
  T{ #introduce { out-d { 6256069 } } }
  T{ #declare { declaration { { 6256069 fixnum } } } }
  T{ #return { in-d V{ 6256069 } } }
}
```

Listing 21: [{ **fixnum** } declare] build-tree

```
V{
  T{ #push { literal "Error!" } { out-d { 6256051 } } }
  T{ #call
    { word throw }
    { in-d V{ 6256051 } }
    { out-d { } }
  }
  T{ #terminate { in-d V{ } } { in-r V{ } } }
  T{ #return { in-d V{ } } }
}
```

Listing 22: ["Error!" **throw**] build-tree

purposefully want to throw an error. **#terminate** is introduced where the program halts prematurely. When checking the stack height, it gets to be treated specially so that *terminated* stack effects unify with any other effect. That way, branches will still be balanced even if one of them unconditionally throws an error. Listing 22 shows **#terminate** being introduced by the **throw** word.

Next, Listing 18 on page 20 defines nodes for branching based off the superclass **#branch**. The **children** slot contains vectors of nodes representing different branches. **live-branches** is filled in during later analyses to indicate which branches are alive so that dead ones may be removed. For instance, **#if** will have two elements in its **children** slot representing the true and false branches. On the other hand, **#dispatch** has an arbitrary number of children. It corresponds to the **dispatch** primitive, which is an implementation detail of the generic word system used to speed up method dispatch.

You may have noted the emphasis on introducing new values, instead of reassigning old ones. Even **#shuffles** output fresh identifiers, letting their values be determined by the **mapping**. The reason for this is that **compiler.tree** uses static single assignment (SSA) form, wherein every variable is defined by exactly one statement. This simplifies the properties of variables, which helps optimizations perform faster and with better results. By giving unique names to the targets of each assignment, the SSA property is guaranteed. However, **#branches** introduce ambiguity: after, say, an **#if**, what will the **out-d** be? It depends on which branch is taken. To remedy this problem, after any **#branch** node, Factor will place a **#phi** node—the classical SSA “phony function”, ϕ . While it doesn’t perform any literal computation, conceptually ϕ selects between its inputs, choosing the “correct” argument depending on control flow. This can then be assigned to

cite?

a unique value, preserving the SSA property. In Factor, this is represented by a `phi-in-d` slot, which is a sequence of sequences. Each element corresponds to the `out-d` of the child at the same position in the `children` of the preceding `#branch` node. The `#phi`'s `out-d` gives unique names to the output values.

```
V{
  T{ #introduce { out-d { 6256247 } } }
  T{ #if
    { in-d { 6256247 } }
    { children
      {
        V{
          T{ #push
            { literal 1 }
            { out-d { 6256248 } }
          }
        }
        V{
          T{ #push
            { literal 2 }
            { out-d { 6256249 } }
          }
        }
      }
    }
  }
  T{ #phi
    { phi-in-d { { 6256248 } { 6256249 } } }
    { out-d { 6256250 } }
    { terminated V{ f f } }
  }
  T{ #return { in-d V{ 6256250 } } }
}
```

Listing 23: [[1] [2] if] build-tree

For example, the `#phi` in Listing 23 will select between the 6256248 return value of the first child or the 6256249 output of the second. Either way, we can refer to the result as 6256250 afterwards. The `terminated` slot of the `#phi` tells us if there was a `#terminate` in any of the branches.

The `#recursive` node encapsulates *inline recursive* words. In Factor, words may be annotated with simple compiler declarations, which guide optimizations. If we follow a standard colon definition with the `inline` word, we're saying that its definition can be spliced into the call-site, rather than generating code to jump to a subroutine. Inline words that call themselves must additionally be declared `recursive`. For example, we could write: `foo (--) foo ; inline recursive.`

The nodes `#enter-recursive`, `#call-recursive`, and `#return-recursive` denote different stages of the recursion—the beginning, recursive call, and end, respectively. They carry around a lot of metadata about the nature of the recursion, but it doesn’t serve our purposes to get into the details. Similarly, we gloss over the final nodes of Listing 18 on page 20 correspond to Factor’s foreign function interface (FFI) vocabulary, called `alien`. At a high level, `#alien-node`, `#alien-invoke`, `#alien-indirect`, `#alien-assembly`, and `#alien-callback` are used to make calls to C libraries from within Factor.

```
: optimize-tree ( nodes -- nodes' )
[
  analyze-recursive
  normalize
  propagate
  cleanup
  dup run-escape-analysis? [
    escape-analysis
    unbox-tuples
  ] when
  apply-identities
  compute-def-use
  remove-dead-code
  ?check
  compute-def-use
  optimize-modular-arithmetic
  finalize
] with-scope ;
```

Listing 24: Optimization passes on the high-level IR

Shouldn’t bold “cleanup” in Listing 24

Now that we’re familiar with the structure of the high-level IR, we can turn our attention to optimization. Listing 24 shows the passes performed on a sequence of nodes by the word `optimize-tree`. Before optimization can begin, we must gather some information and clean up some oddities in the output of `build-tree`. `analyze-recursive` is called first to identify and mark loops in the tree. Effectively, this means we detect tail-recursion introduced by `#recursive` nodes. Future passes can then use this information for data flow analysis. Then, `normalize` makes the tree more consistent by doing two things:

- All `#introduce` nodes are removed and replaced by a single `#introduce` at the beginning of the whole program. This way, further passes needn’t handle `#introduce` nodes.
- As constructed, the `in-d` of a `#call-recursive` will be the entire stack at the time of the call. This assumption happens because we don’t know how many inputs it needs until the `#return-recursive` is processed, because of row polymorphism. So, here we figure out

exactly what stack entries are needed, and trim the `in-d` and `out-d` of each `#call-recursive` accordingly.

Once these passes have cleaned up the tree, `propagate` performs probably the most extensive analysis of all the phases. In short, it performs an extended version of sparse conditional constant propagation (SCCP). The traditional data flow analysis combines global copy propagation, constant propagation, and constant folding in a *flow-sensitive* way. That is, it will propagate information from branches that it knows are definitely taken (e.g., because `#if` is always given a true input). Instead of using the typical single-level (numeric) constant value lattice, Factor uses a lattice augmented by information about classes, numeric value ranges, array lengths, and tuple slots' classes. Classes can be used in the lattice with the partial-order protocol described briefly in Section 2.4. Additionally, the transfer functions are allowed to inline certain calls if enough information is present. This occurs in the transfer function since generic words' inline expansions into particular methods provide more information, thus giving us more opportunities for propagation. This is particularly useful for arithmetic words. In Factor, words like `+` and `*` are generics that work across all sorts of numeric representations, be they `fixnums`, `floats`, `bignums`, etc. If the operation overflows, the values are automatically cast up to larger representations. But iterated refinement of the inputs' classes can let the compiler select more specific, efficient methods (e.g., if both arguments are `fixnums`).

cite

Interval propagation also helps propagate class information. By refining the range of possible values a particular item can have, we might discover that, say, it's small enough to fit in a `fixnum` rather than a `bignum`. There are plenty more things that interval propagation can tell us, too. For example, it may give us enough information to remove overflow checks performed by numeric words. And if the interval has zero length, we may replace the value with a constant. This then continues getting propagated, contributing to constant folding and so forth.

`propagate` iterates through the nodes collecting all of this data until reaching a stable point where inferences can no longer be drawn. Technically, this information doesn't alter the tree at all; we merely store it so that speculative decisions may be realized later. The next word in Listing 24 on the previous page, `cleanup`, does just this by inlining words, folding constants, removing overflow checks, deleting unreachable branches, and flattening inline-recursive words that don't actually wind up calling themselves (e.g., because the calls got constant-folded).

The next major pass is `escape-analysis`, whose information is used for the actual transformation `unbox-tuples`. This discovers tuples that *escape* by getting passed outside of a word. For instance, the inputs to `#return` obviously escape, as they are passed to the world outside of the word in question. Similarly, inputs to the `#call` of another word escape. So, though the tuples in `escaping-via-#return` and `escaping-via-#call` in Listing 25 on the following page both escape, we can see the one in `non-escaping` does not. In fact, the last allocation is unnecessary. By identifying this, `unbox-tuples` can then rewrite the code to avoid allocating a `data-struct` altogether, instead manipulating the slots' values directly. Note that this only happens for immutable tuples, all of whose slots are `read-only`. Otherwise, we would need to perform more advanced pointer analyses to discover aliases.

`apply-identities` follows to simplify words with known identity elements. If, say, an argument to `+` is 0, we can simply return the other argument. This converts the `#call` to `+` into a simple `#shuffle`. These identities are defined for most arithmetic words.

```

TUPLE: data-struct
  { a read-only }
  { b read-only } ;

: escaping-via-#return ( -- data-struct )
  1 2 data-struct boa ;

: escaping-via-#call ( -- )
  1 2 data-struct boa pprint ;

: non-escaping ( -- )
  1 2 data-struct [ a>> ] [ b>> ] bi + ;

```

Listing 25: Escaping vs. non-escaping tuple allocations

Another simple few passes come next in Listing 24 on page 24. True to its name, `compute-def-use` computes where SSA values are defined and used. Values that are never used are eliminated by `remove-dead-code`. `?check` conditionally performs some consistency checks on the tree, mostly to make sure that no errors were introduced in the stack flow. If a global variable isn’t toggled on, this part is skipped. We run `compute-def-use` again to update the information after altering the tree with dead code elimination.

Finally, `optimize-modular-arithmetic` performs a form of strength-reduction on arithmetic words that only use the low-order bits of their inputs/results, which may also remove more unnecessary overflow checks. `finalize` cleans up a few random miscellaneous bits of the tree (removing empty shufflers, deleting `#copy` nodes, etc.) in preparation for lower-level optimizations.

Double-check zealous syntax-highlighting

3.3 Low-level Optimizations

The low-level IR in `compiler.cfg` takes the more conventional form of a control flow graph (CFG). A CFG (not to be confused with “context-free grammar”) is an arrangement of instructions into *basic blocks*: maximal sequences of “straight-line” code, where control does not transfer out of or into the middle of the block. Directed edges are added between blocks to represent control flow—either from a branching instruction to its target, or from the end of a basic block to the start of the next one. Construction of the low-level IR proceeds by analyzing the control flow of the high-level IR and converting the nodes of Section 3.2 into lower-level, more conventional instructions modeled after typical assembly code. There are over a hundred of these instructions, but many are simply different versions of the same operation. For instance, while instructions are generally called on *virtual registers* (represented in Factor simply by integers), there are *immediate* versions of instructions. The `##add` instruction, as an example, represents the sum of the contents of two registers, but `##add-imm` sums the contents of one register and an integer literal. Other instructions are inserted to make stack reads and writes explicit, as well as to balance the height. Below is a categorized list of all the instruction objects (each one is a subclass of the `insn` tuple).

cite

Is the complete list really necessary?

- Loading constants: `##load-integer`, `##load-reference`
- Optimized loading of constants, inserted by representation selection: `##load-tagged`, `##load-float`, `##load-double`, `##load-vector`
- Stack operations: `##peek`, `##replace`, `##replace-imm`, `##inc-d`, `##inc-r`
- Subroutine calls: `##call`, `##jump`, `##prologue`, `##epilogue`, `##return`
- Inhibiting tail-call optimization (TCO): `##no-tco`
- Jump tables: `##dispatch`
- Slot access: `##slot`, `##slot-imm`, `##set-slot`, `##set-slot-imm`
- Register transfers: `##copy`, `##tagged>integer`
- Integer arithmetic: `##add`, `##add-imm`, `##sub`, `##sub-imm`, `##mul`, `##mul-imm`, `##and`, `##and-imm`, `##or`, `##or-imm`, `##xor`, `##xor-imm`, `##shl`, `##shl-imm`, `##shr`, `##shr-imm`, `##sar`, `##sar-imm`, `##min`, `##max`, `##not`, `##neg`, `##log2`, `##bit-count`
- Float arithmetic: `##add-float`, `##sub-float`, `##mul-float`, `##div-float`, `##min-float`, `##max-float`, `##sqrt`
- Single/double float conversion: `##single>double-float`, `##double>single-float`
- Float/integer conversion: `##float>integer`, `##integer>float`
- SIMD operations: `##zero-vector`, `##fill-vector`, `##gather-vector-2`, `##gather-int-vector-2`, `##gather-vector-4`, `##gather-int-vector-4`, `##select-vector`, `##shuffle-vector`, `##shuffle-vector-halves-imm`, `##shuffle-vector-imm`, `##tail>head-vector`, `##merge-vector-head`, `##merge-vector-tail`, `##float-pack-vector`, `##signed-pack-vector`, `##unsigned-pack-vector`, `##unpack-vector-head`, `##unpack-vector-tail`, `##integer>float-vector`, `##float>integer-vector`, `##compare-vector`, `##test-vector`, `##test-vector-branch`, `##add-vector`, `##saturated-add-vector`, `##add-sub-vector`, `##sub-vector`, `##saturated-sub-vector`, `##mul-vector`, `##mul-high-vector`, `##mul-horizontal-add-vector`, `##saturated-mul-vector`, `##div-vector`, `##min-vector`, `##max-vector`, `##avg-vector`, `##dot-vector`, `##sad-vector`, `##horizontal-add-vector`, `##horizontal-sub-vector`, `##horizontal-shl-vector-imm`, `##horizontal-shr-vector-imm`, `##abs-vector`, `##sqrt-vector`, `##and-vector`, `##andn-vector`, `##or-vector`, `##xor-vector`, `##not-vector`, `##shl-vector-imm`, `##shr-vector-imm`, `##shl-vector`, `##shr-vector`
- Scalar/vector conversion: `##scalar>integer`, `##integer>scalar`, `##vector>scalar`, `##scalar>vector`

- Boxing and unboxing aliens: `##box-alien`, `##box-displaced-alien`, `##unbox-any-c-ptr`, `##unbox-alien`
- Zero-extending and sign-extending integers: `##convert-integer`
- Raw memory access: `##load-memory`, `##load-memory-imm`, `##store-memory`, `##store-memory-imm`
- Memory allocation: `##allot`, `##write-barrier`, `##write-barrier-imm`, `##alien-global`, `##vm-field`, `##set-vm-field`
- The FFI: `##unbox`, `##unbox-long-long`, `##local-allot`, `##box`, `##box-long-long`, `##alien-invoke`, `##alien-indirect`, `##alien-assembly`, `##callback-inputs`, `##callback-outputs`
- Control flow: `##phi`, `##branch`
- Tagged conditionals: `##compare-branch`, `##compare-imm-branch`, `##compare`, `##compare-imm`
- Integer conditionals: `##compare-integer-branch`, `##compare-integer-imm-branch`, `##test-branch`, `##test-imm-branch`, `##compare-integer`, `##compare-integer-imm`, `##test`, `##test-imm`
- Float conditionals: `##compare-float-ordered-branch`, `##compare-float-unordered-branch`, `##compare-float-ordered`, `##compare-float-unordered`
- Overflowing arithmetic: `##fixnum-add`, `##fixnum-sub`, `##fixnum-mul`
- GC checks: `##save-context`, `##check-nursery-branch`, `##call-gc`
- Spills and reloads, inserted by the register allocator: `##spill`, `##reload`

```

: optimize-cfg ( cfg -- cfg' )
  optimize-tail-calls
  delete-useless-conditionals
  split-branches
  join-blocks
  normalize-height
  construct-ssa
  alias-analysis
  value-numbering
  copy-propagation
  eliminate-dead-code ;

```

Listing 26: Optimization passes on the low-level IR

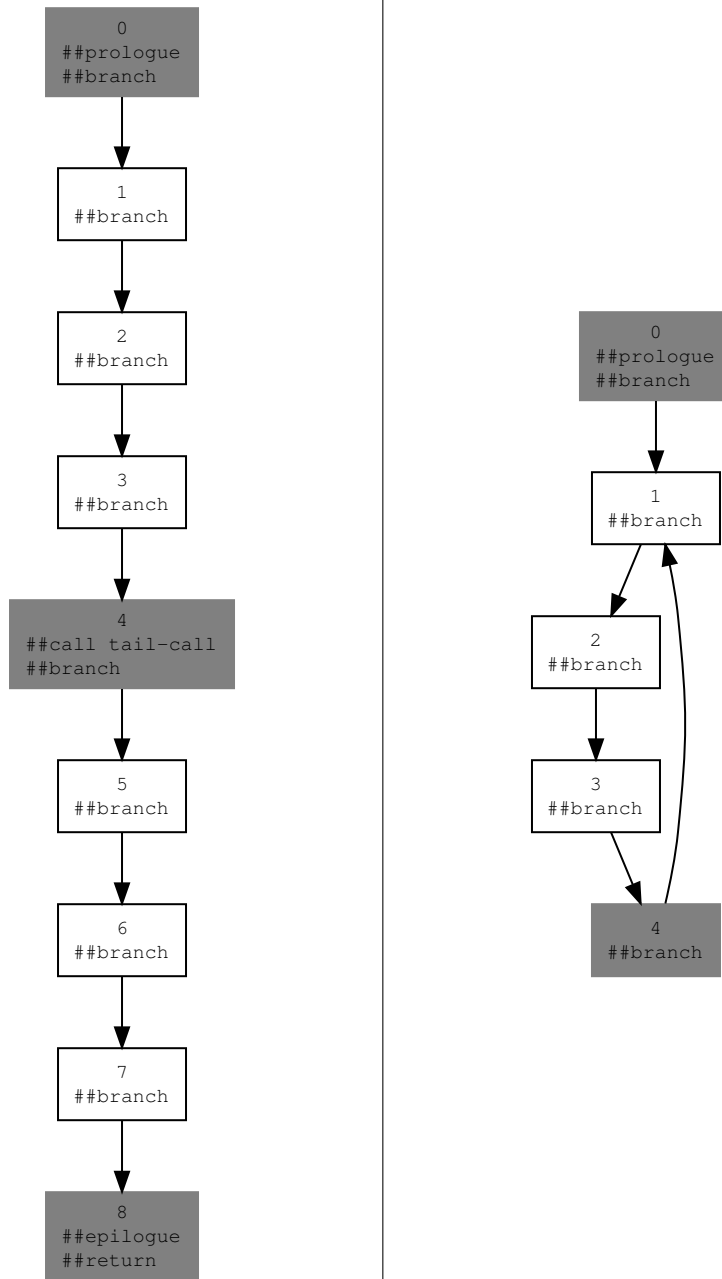


Figure 5: `tail-call` before and after `optimize-tail-calls`

By translating the high-level IR into instructions that manipulate registers directly, we reveal further redundancies that can be optimized away. The `optimize-cfg` word in Listing 26 on page 28 shows the passes performed in doing this. The first word, `optimize-tail-calls`, performs tail call elimination on the CFG. *Tail calls* are those that occur within a procedure and whose results are immediately returned by that procedure. Instead of allocating a new call stack frame, we may convert tail calls into simple jumps, since afterwards the current procedure’s call frame isn’t really needed. In the case of recursive tail calls, we can convert special cases of recursion into loops in the CFG, so that we won’t trigger call stack overflows. For instance, consider Figure 5 on the previous page, which shows the effect of `optimize-tail-calls` on the following definition:

```
: tail-call ( -- ) tail-call ;
```

Note the recursive call (trivially) occurs at the end of the definition, just before the return point. When translated to a CFG, this is a `##call` instruction, as seen in block 4 to the left of Figure 5 on the preceding page. This is also just before the final `##epilogue` and `##return` instructions in block 8, as blocks 5–7 are effectively empty (these excessive `##branches` will be eliminated in a later pass). Because of this, rather than make a whole new subroutine call, we can convert it into a `##branch` back to the beginning of the word, as in the CFG to the right.

The next pass in Listing 26 on page 28 is `delete-useless-conditionals`, which removes branches that go to the same basic block. This situation might occur as a result of optimizations performed in the high-level IR. To see it in action, Figure 6 on the following page shows the transformation on a purposefully useless conditional, `[] [] if`. Before removing the useless conditional, the CFG `##peek`s at the top of the data stack (`D 0`), storing the result in the virtual register 1. This value is popped, so we decrement the stack height (`##inc-d -1`). Then, `##compare-imm-branch` in block 2 compares the value in the virtual register 2 (which is a copy of 1, the top of the stack) to the immediate value `f` to see if it’s not equal (signified by `cc/=`). However, both branches jump through several empty blocks and merge at the same destination. Thus, we can remove both branches and replace `##compare-imm-branch` with an unconditional `##branch` to the eventual destination. We see this on the right of Figure 6 on the next page.

In order to expose more opportunities for optimization, `split-branches` will actually duplicate code. We use the fact that code immediately following a conditional will be executed along either branch. If it’s sufficiently short, we copy it up into the branches individually. That is, we change `[A] [B] if C` into `[A C] [B C] if`, as long as `C` is small enough. Later analyses may then, for example, more readily eliminate one of the branches if it’s never taken. Figure 7 on page 32 shows what such a transformation looks like on a CFG. The example `[1] [2] if dup` is essentially changed into `[1 dup] [2 dup] if`, thus splitting the block with two predecessors (block 9) on the left.

The next pass, `join-blocks`, compacts the CFG by joining together blocks involved in only a single control flow edge. Mostly, this is to clean up the myriad of empty or short blocks introduced during construction, like sequences of a bunch of `##branches`. Figure 8 on page 33 shows this pass on the CFG of `0 100 [1 fixnum+fast] times`. `fixnum+fast` is a specialized version of `+` that suppresses overflow and type checks. We use it here to keep the CFG simple. We’ll be using this particular code to illustrate all but one of the remaining optimization passes in Listing 26 on page 28, as it’s a motivating example for the work in this thesis. The passes before `join-blocks`

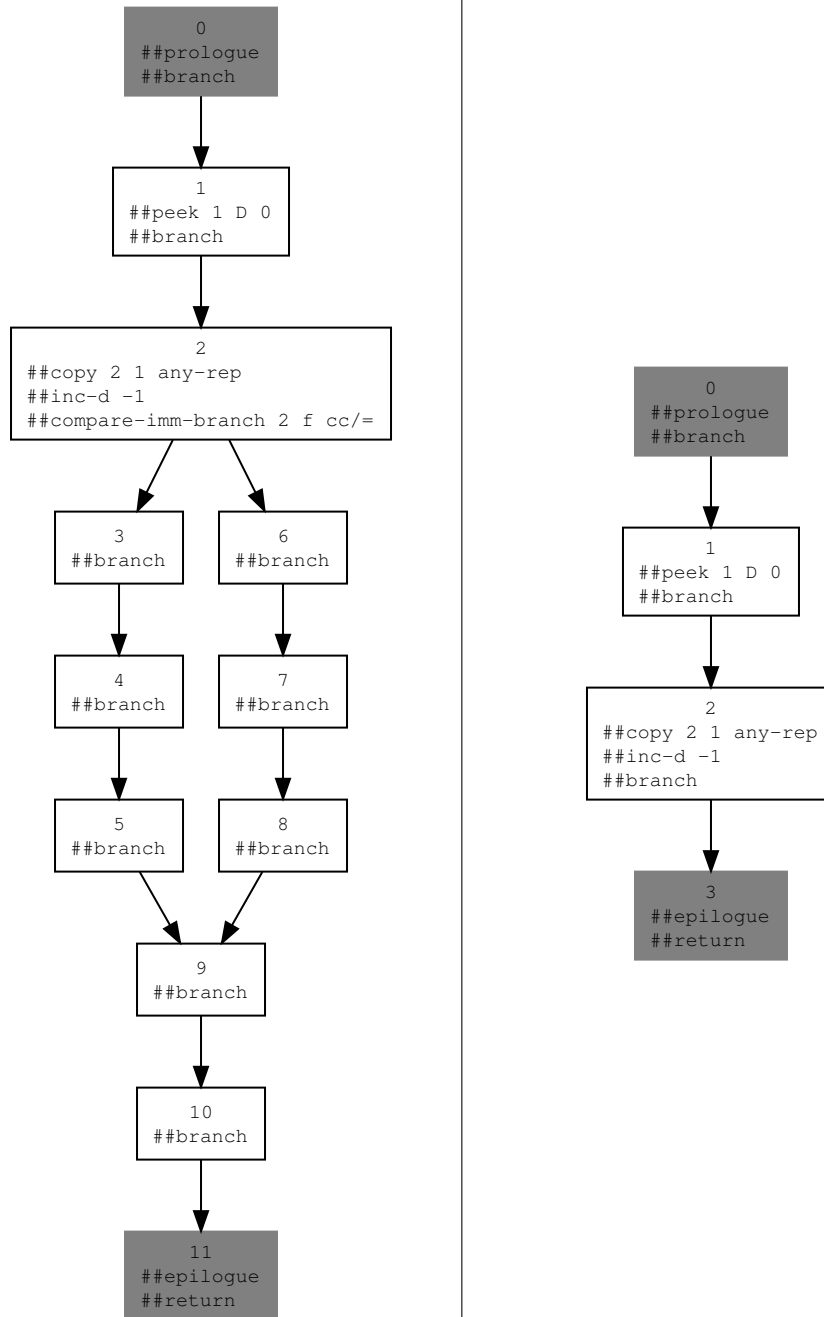


Figure 6: [] [] `if` before and after `delete-useless-conditionals`

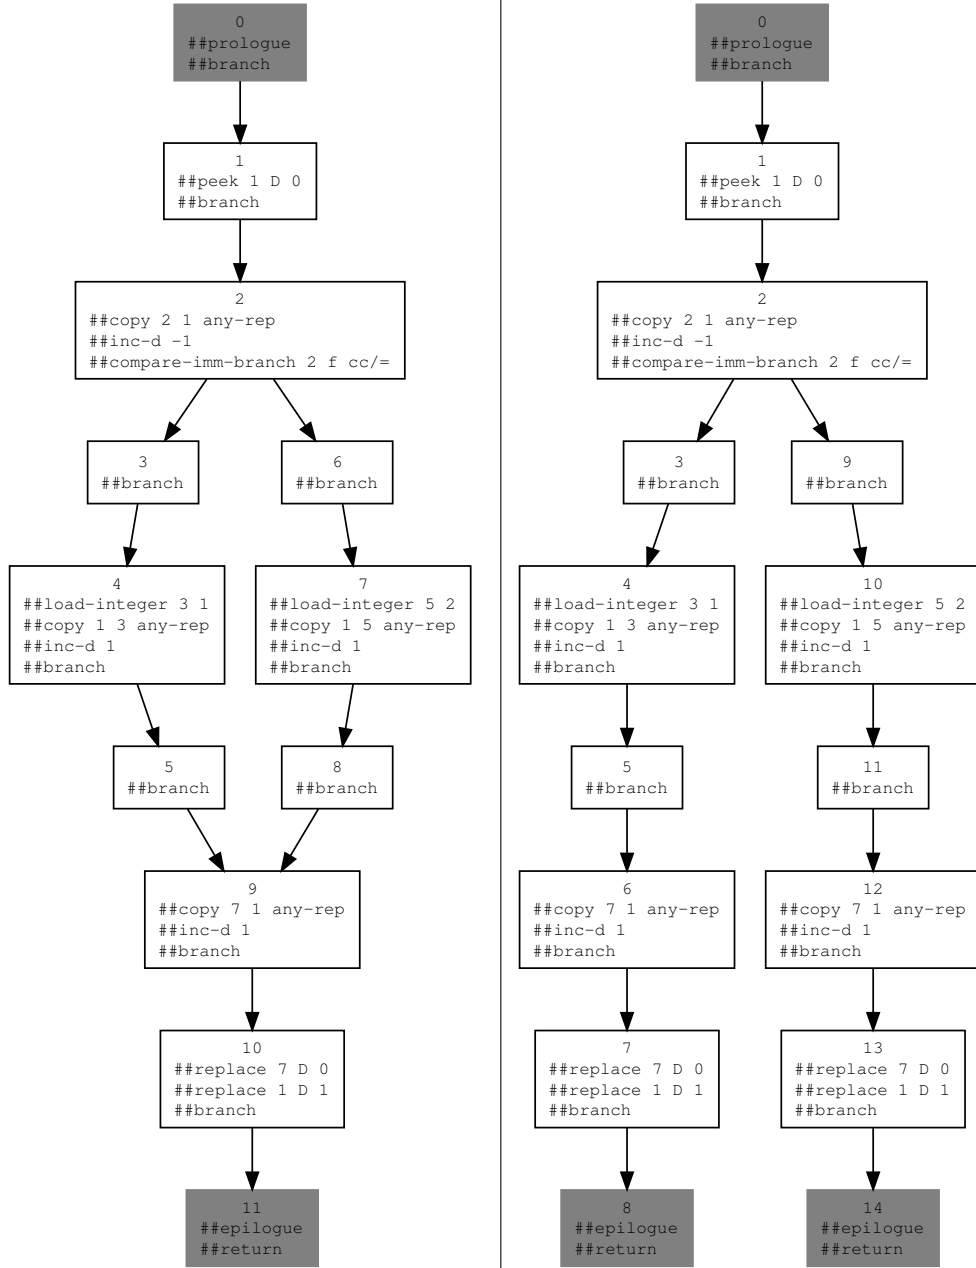


Figure 7: [1] [2] if dup before and after split-branches

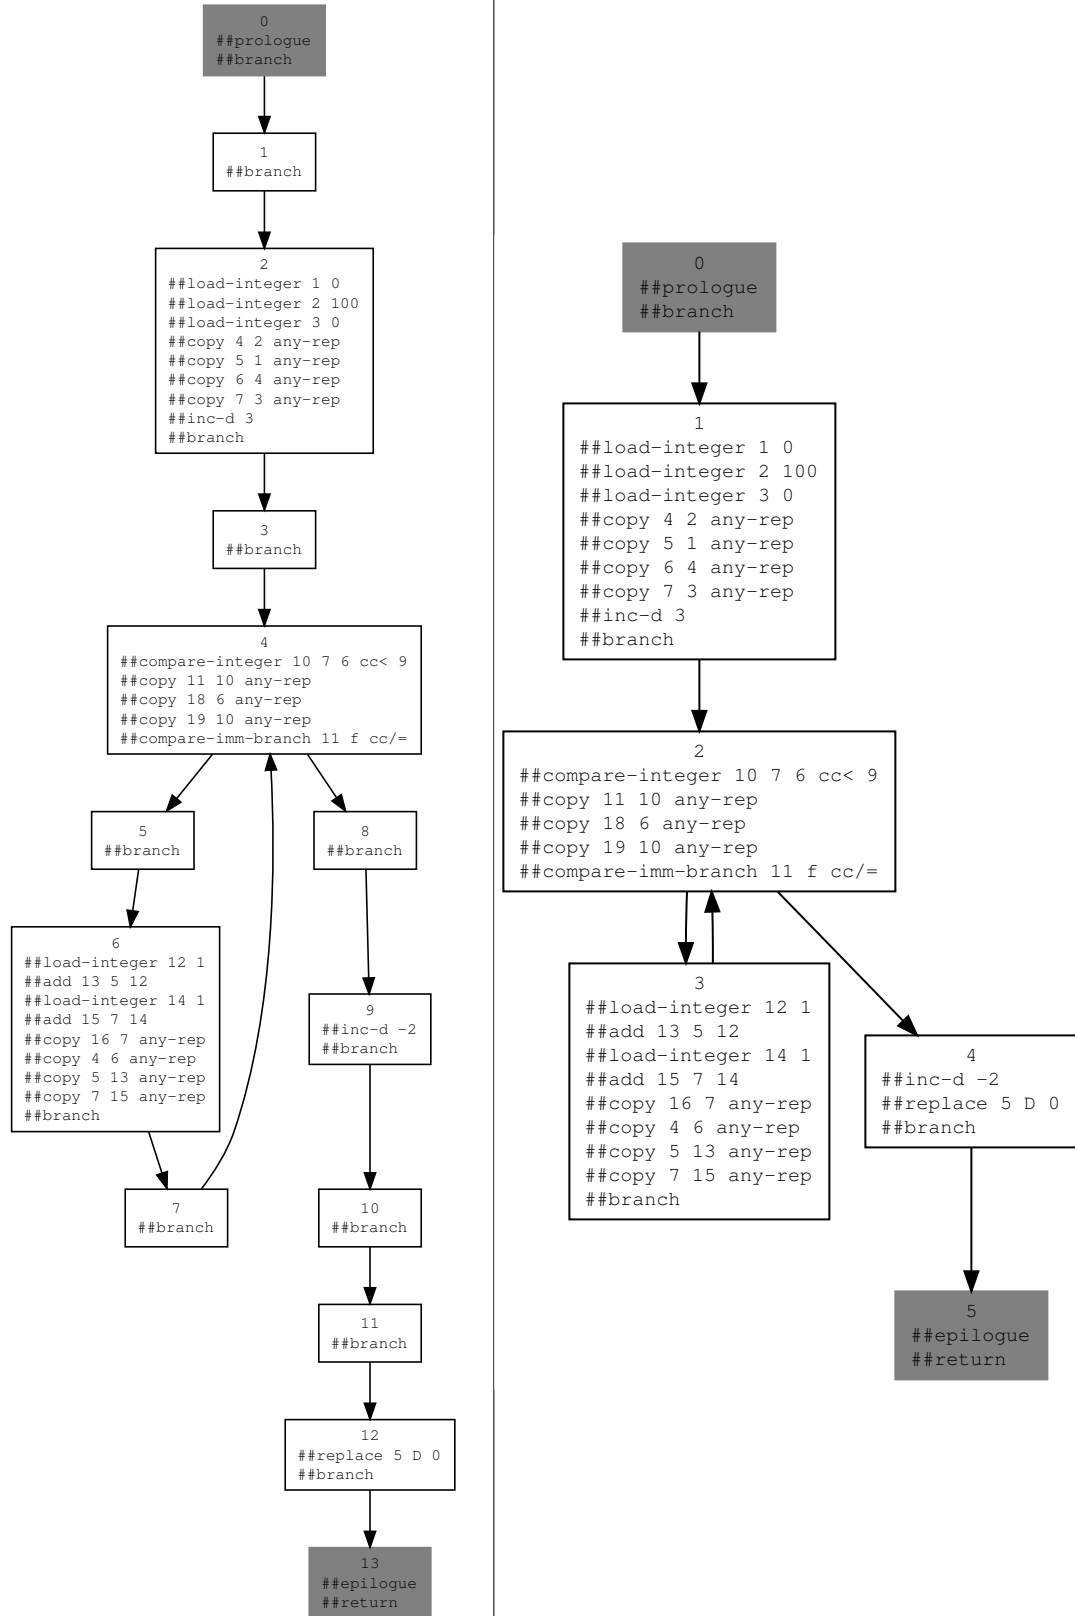


Figure 8: 0 100 [1 fixnum+fast] times before and after join-blocks

don't change the CFG seen on the left in Figure 8 on the previous page, but we get rid of the useless `##branch` blocks in the CFG on the right.

Figure 9 on the following page shows the result of applying `normalize-height` to the result of `join-blocks`. This phase combines and canonicalizes the instructions that track the stack height, like `##inc-d`. While the shuffling in this example isn't complex enough to be interesting, neither is this phase. It amounts to more cleanup: multiple height changes are combined into single ones at the beginnings of the basic blocks. In Figure 9 on the next page, this means that `##inc-d` is moved to the top of block 1, as compared to the right of Figure 8 on the preceding page.

In converting the high-level IR to the low-level, we actually lose the SSA form of `compiler.tree`. Not only does the construction do this, but `split-branches` also copies basic blocks verbatim, so any value defined will have a duplicate definition site, violating the SSA property. `construct-ssa` recomputes a so-called *pruned* SSA form, wherein ϕ functions are inserted only if the variables are live after the insertion point. This cuts down on useless ϕ functions. Figure 10 on page 36 shows the reconstructed SSA form of the CFG from Figure 9 on the following page.

cite
TDMSC
and con-
struction
algorithm

The next pass, `alias-analysis`, doesn't change the CFG of 0 100 [1 fixnum+fast] times, so we won't have an accompanying figure. At a high level, `alias-analysis` is easy to understand: it eliminates redundant memory loads and stores by rewriting certain patterns of memory access. If the same location is loaded after being stored, we convert the latter load into a `##copy` of the value we stored. Two reads of the same location with no intermittent write gets the second read turned into a `##copy`. Similarly, if we see two writes without a read in the middle, the first write can be removed.

`value-numbering` is the key focus of this thesis. It will be detailed in Section 4 on page 41. For now, it does to think of it as a combination of common subexpression elimination and constant folding. In Figure 11 on page 37, we see several changes:

- `##load-integer 23 0` in block 1 of Figure 10 on page 36 (which assigns the value 0 to the virtual register 23) is redundant, so is replaced by `##copy 23 21`.
- In block 2, the last instruction `##compare-imm-branch 32 f cc/=` is the same as `##compare-integer-branch 30 26 cc<`. The source register (32) of the original is a `##copy` of 31, which itself is computed by `##compare-integer 31 30 26 cc< 9`. So, the `##compare-imm-branch` is equivalent to a simple `##compare-integer-branch`, which doesn't use the temporary virtual register 9 and doesn't waste time comparing against the `f` object.
- The second operands in both `##adds` of block 3 are just constants stored by `##load-integers`. So, these are turned into `##add-imms`.
- Also, the second `##load-integer` in block 3 just loads 1 like the first instruction. Therefore, it's replaced by a `##copy`.

In Section 4, we'll see how and why this pass fails to identify other equivalences.

Following `value-numbering`, `copy-propagation` performs a global pass that eliminates `##copy` instructions. Uses of the copies are replaced by the originals. So, in Figure 12 on page 38, we can

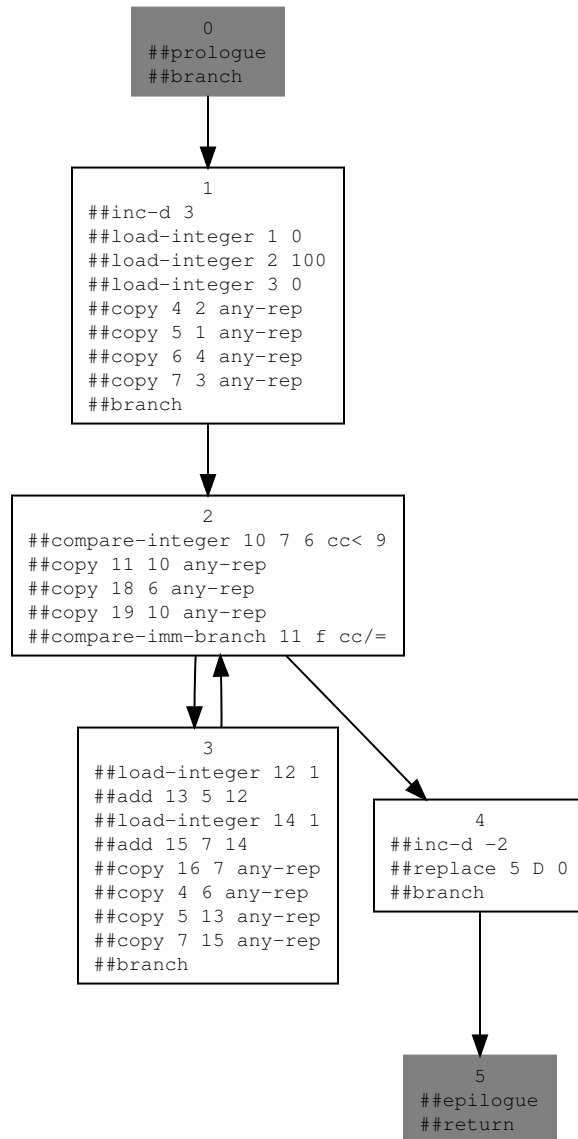


Figure 9: 0 100 [1 fixnum+fast] **times** after normalize-height




Figure 10: 0 100 [1 fixnum+fast] **times** after construct-ssa

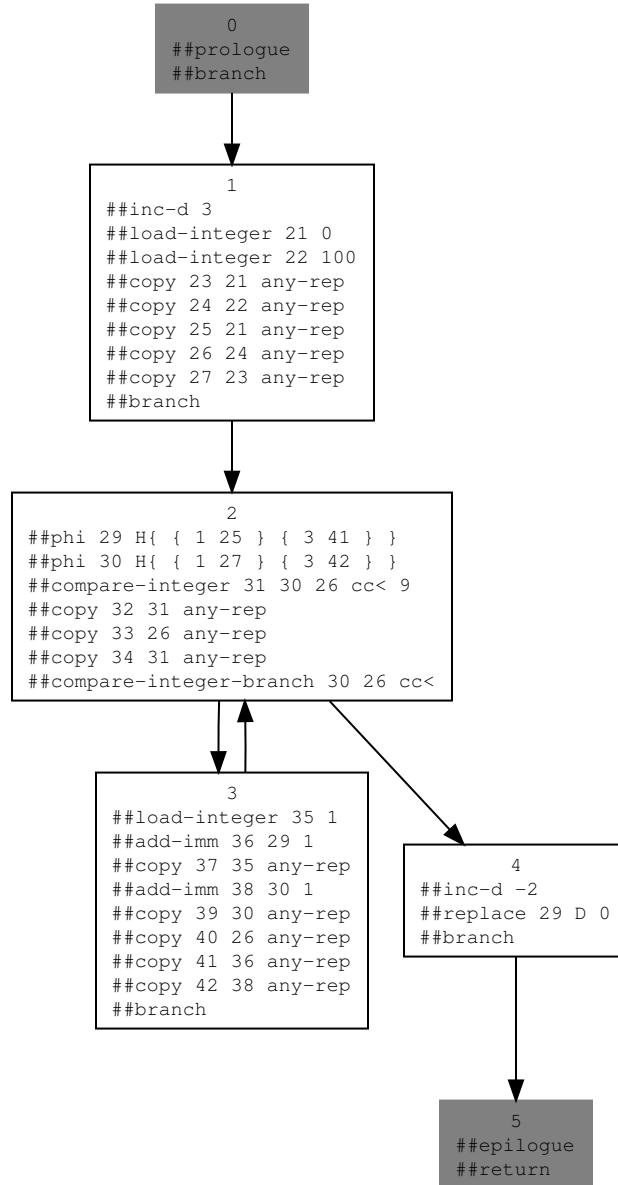


Figure 11: 0 100 [1 fixnum+fast] **times** after value-numbering

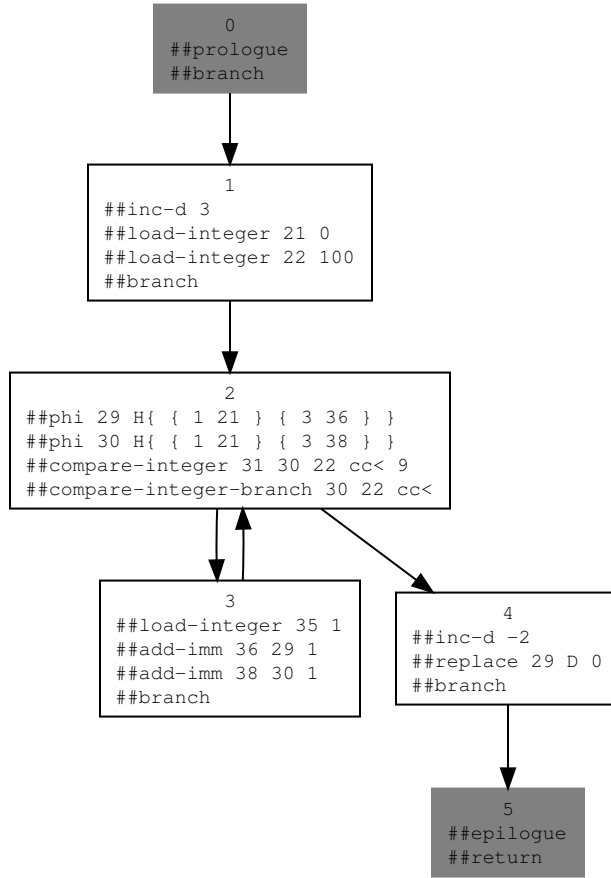


Figure 12: 0 100 [1 fixnum+fast] **times** after copy-propagation




Figure 13: 0 100 [1 fixnum+fast] **times** after eliminate-dead-code

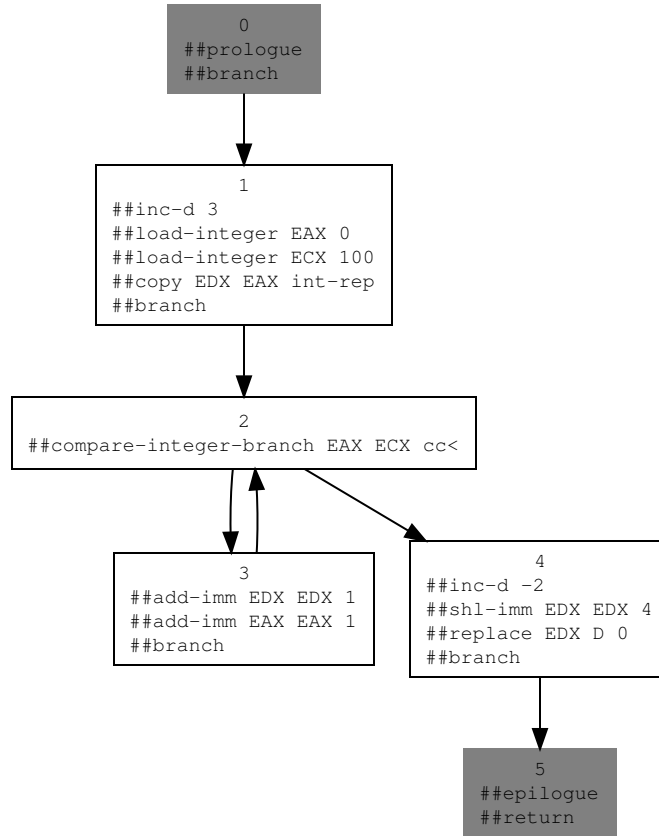


Figure 14: 0 100 [1 fixnum+fast] times after finalize-cfg

see that all of the `##copy` instructions have been removed and, for instance, the use of the virtual register 25 in block 2 has been replaced by 21, since 25 was a copy of it.

Next, dead code is removed by `eliminate-dead-code`. Figure 13 on the previous page shows that the `##compare-integer` in block 2 and the `##load-integer` in block 3 were removed, since they defined values that were never used.

The final pass in Listing 26 on page 28, `finalize-cfg`, itself consists of several more passes. We will not get into many details here, but at a high level, the most important passes figure out how virtual registers should map to machine registers. We first figure out when certain values can be unboxed. Then, instructions are reordered in order to reduce *register pressure*. That is, we try to schedule instructions around each other so that we don't need to store more values than we have machine registers. That way, we avoid *spilling* registers onto the heap, which wastes time. After leaving SSA form, we perform a *linear scan* register allocation, which replaces virtual registers with machine registers and inserts `##spill` and `##reload` instructions for the cases we can't avoid. Figure 14 shows an example on an Intel x86 machine, which has enough registers that we needn't spill anything.

4 Value Numbering

At a very basic level, most optimization techniques revolve around avoiding redundant or unnecessary computation. Thus, it's vital that we discover which values in a program are equal. That way, we can simplify the code that wastes machine cycles repeatedly calculating the same values. Classic optimization phases like constant/copy propagation, common subexpression elimination, loop-invariant code motion, induction variable elimination, and others discussed in the de facto treatise, “The Dragon Book”, perform this sort of redundancy elimination based on information about the equality of expressions.

cite

In general, the problem of determining whether two expressions in a program are equivalent is undecidable. Therefore, we seek a *conservative* solution that doesn't necessarily identify all equivalences, but is nevertheless correct about any equivalences it does identify. Solving this equivalence problem is the work of *value numbering* algorithms. These assign every value in the program a number such that two values have the same value number if and only if the compiler can prove they will be equal at runtime.

Value numbering has a long history in literature and practice, spanning many techniques. In Section 3.3 we saw the **value-numbering** word, which is actually based on some of the earliest—and least effective—methods of value numbering. Section 4.1 describes the way Factor's current algorithm works, and highlights its shortcomings to motivate the main work of this thesis, which is covered in Sections 4.2 and 4.3. We finish the section by analyzing the results of these changes and reviewing the literature for further enhancements that can be made to this optimization pass.

4.1 Local Value Numbering

Tracing the exact origins of value numbering is difficult. It's thought to have originally been invented in the 1960s by Balke. The earliest tangible reference to a value numbering (at least, the earliest point where discussions in the literature seem to start) appears in an oft-cited but unpublished work of Cocke. The technique is relatively simple, but not as powerful as other methods for reasons described hereafter.

cite Simpson

cite-like

The algorithm considers a single basic block. For each instruction (from top to bottom) in the block, we essentially let the value number of the assignment target be a hash of the operator and the value numbers of the operands. That is, we hash the *expression* being computed by an instruction. Thus, assuming a proper hash function, two expressions are *congruent* if

- they have the same operators and
- their operands are congruent.

This is our approximation of runtime equivalence. The first property is fulfilled by basing the hash, in part, on the operator. The second property holds because the hash is based on the value numbers of the statement's operands—not just the operands as they appear in code (i.e., *lexical* equivalence). Any information about congruence is propagated through the value numbers. We'll have discovered any such equivalences among the operands before computing the value number of the assignment target because every value in a basic block is either defined before it's used, or

else defined at some point in a predecessor of the block, which we don't care about when only considering one basic block.

This is the first shortcoming of the algorithm. It is *local*, focusing on only one basic block at a time. Any definitions outside the boundaries of the basic block won't be reused, even if they reach the block. This severely limits the scope of the redundancies we can discover. We could improve upon this by considering the algorithm across an entire loop-free CFG in any *topological order*. In such an ordering, a basic block B comes before any other block B' to which it has an edge. Thus, any “outside” variables that instructions in B' rely on must have come from B or earlier, which will have already been computed in a traversal of such an ordering. However, CFGs usually contain cycles or loops (at least interesting ones do), which make such an ordering impossible. We could still pick a topological order that ignores back-edges, but we may encounter operands whose values flow along those back-edges, so haven't been processed yet. We'll address this issue later.

In Factor, expressions are basically instructions (the `insn` objects discussed in Section 3.3) that have had their destination registers stripped. Instructions can be converted to expressions with the `>expr` word defined in the `compiler.cfg.value-numbering.expressions` vocabulary. For instance, an `##add` instruction with the destination register 1 and source registers 2 and 3 will be converted into an array of three elements:

- The `##add` class word, indicating the expression is derived from an `##add` instruction.
- The value number of the virtual register 2.
- The value number of the virtual register 3.

Some instructions are not *referentially transparent*, meaning they can't be replaced with the value they compute without changing the program's behavior. For example, `##call` and `##branch` cannot reasonably be converted into expressions. In these cases, `>expr` merely returns a unique value.

The hashing of expressions takes place in the so-called *expression graph* implemented in the vocabulary shown in Listing 27 on the following page. This consists of three global hash tables that relate virtual registers, value numbers, instructions, and expressions. Since virtual registers are just integers, we actually use them as value numbers, too. `vregs>vns` maps virtual registers to their value numbers. If a virtual register is mapped to itself in this table, its definition is the canonical instruction that we use to compute the value. This instruction is stored in the `vns>insns` table. Finally, the most important mapping is `exprs>vns`. True to its name, it uses expressions as keys, which of course are implicitly hashed. Thus, we can use this table to determine equivalence of expressions.

Other definitions in Listing 27 on the next page manipulate expressions and the graph. The global variable `input-expr-counter` is used in the generation of unique expressions discussed earlier. `init-value-graph` initializes this and all the tables. `set-vn` establishes a mapping from a virtual register to a value number in `vregs>vns`. `vn>insn` gives terse access to the `vns>insns` table. `vreg>insn` uses `vregs>vns` and `vns>insns` to get the canonical instruction that defines a given virtual register. Finally, `vreg>vn` looks up the value of a key in the `vregs>vns` table. Importantly, if the key is not yet present in the table, it is automatically mapped to itself—it's assumed that the virtual register does not correspond to a redundant instruction.

```

! Copyright (C) 2008, 2010 Slava Pestov.
! See http://factorcode.org/license.txt for BSD license.
USING: accessors kernel math namespaces assocs ;
IN: compiler.cfg.value-numbering.graph

SYMBOL: input-expr-counter

! assoc mapping vregs to value numbers
! this is the identity on canonical representatives
SYMBOL: vregs>vns

! assoc mapping expressions to value numbers
SYMBOL: exprs>vns

! assoc mapping value numbers to instructions
SYMBOL: vns>insns

: vn>insn ( vn -- insn ) vns>insns get at ;

: vreg>vn ( vreg -- vn ) vregs>vns get [ ] cache ;

: set-vn ( vn vreg -- ) vregs>vns get set-at ;

: vreg>insn ( vreg -- insn ) vreg>vn vn>insn ;

: init-value-graph ( -- )
  0 input-expr-counter set
  H{ } clone vregs>vns set
  H{ } clone exprs>vns set
  H{ } clone vns>insns set ;

```

Listing 27: The compiler.cfg.value-numbering.graph vocabulary

This is the second shortcoming of the algorithm. It must make a *pessimistic* assumption about congruences. That is, it starts by assuming that every expression has a unique value number, then tries to prove that there are some values which are actually congruent. This fails to discover congruences for values that flow along back-edges, whether we consider a single basic block or an entire topological ordering.

On the other hand, an advantage of this local value numbering algorithm is its simplicity. It makes a single pass over all the instructions, identifying and replacing redundancies *online* (i.e., rewriting as it goes). It's straightforward to write, and even to extend. In particular, there's nothing stopping the online replacements from being more complex than `##copy` instructions. At every step, the currently known value numbers will be sound, and we can use this information for copy/constant propagation, constant folding, and common subexpression elimination.

```
: value-numbering-step ( insns -- insns' )
  init-value-graph
  [ process-instruction ] map flatten ;

: value-numbering ( cfg -- cfg )
  dup [ value-numbering-step ] simple-optimization

  cfg-changed predecessors-changed ;
```

Listing 28: Main words from `compiler.cfg.value-numbering`

To see how Factor accomplishes these extensions, we'll take a look at the `compiler.cfg.value-numbering` vocabulary. Listing 28 shows the main words that start the optimization pass. The `value-numbering-step` word is called on the sequence of instructions that comprise each basic block. It starts the expression graph from a blank slate with `init-value-graph`, then maps the word `process-instruction` on each of them. This is a generic word that we'll study momentarily; it returns either a single `insn` object or a sequence of them (in the case that an instruction is replaced by several others). Then, the work of `value-numbering` is to just call `value-numbering-step` on each basic block, which is done with a combinator called `simple-optimization`. The words `cfg-changed` and `predecessors-changed` alter some internal state of the CFG that has been potentially invalidated by some transformations performed by `process-instruction`.

The methods of `process-instruction` are shown in Listing 29 on the next page. The default behavior for dispatching on an `insns` is to invoke yet another generic word, `rewrite`. This word will return either a replacement `insn` (or sequence thereof) or `f`, indicating that no change has taken place. Thus, by recursively calling `process-instruction`, we can do more specialized processing on this rewritten replacement (e.g., dispatching on `insn` again, which applies `rewrite` once more). If the instruction can't be simplified further, we simply return it. (Note that `[X] [Y] ?if` is the same as `dup [nip X] [drop Y] if`.)

line numbers?

For instances of `foldable-insn` (i.e., `insns` that can be converted to useful expressions with `>expr`), we similarly invoke `rewrite` recursively until no more rewriting occurs. When that happens, rather than just return the instruction, we invoke `check-redundancy`—though only if the instruction defines exactly 1 virtual register, which will be stored in a slot named `dst`. `check-redundancy`

```

GENERIC: process-instruction ( insn -- insn' )

: redundant-instruction ( insn vn -- insn' )
  [ dst>> ] dip [ swap set-vn ] [ <copy> ] 2bi ;

:: useful-instruction ( insn expr -- insn' )
  insn dst>> :> vn
  vn vn vregs>vns get set-at
  vn expr exprs>vns get set-at
  insn vn vns>insns get set-at
  insn ;

: check-redundancy ( insn -- insn' )
  dup >expr dup exprs>vns get at
  [ redundant-instruction ] [ useful-instruction ] ?if ;

M: insn process-instruction
  dup rewrite [ process-instruction ] [ ] ?if ;

M: foldable-insn process-instruction
  dup rewrite
  [ process-instruction ]
  [ dup defs-vregs length 1 = [ check-redundancy ] when ] ?if ;

M: ##copy process-instruction
  dup [ src>> vreg>vn ] [ dst>> ] bi set-vn ;

M: array process-instruction
  [ process-instruction ] map ;

```

Listing 29: The workhorse of `compiler.cfg.value-numbering`

checks if the expression being computed by the instruction is already a key of the `exprs>vns` table. If it is, the instruction is redundant, and we call **redundant-instruction**; otherwise, we call **useful-instruction**. The former uses **set-vn** to map the instruction's `dst` virtual register to the same value number as the expression that was a key of `exprs>vns`. Since value numbers are actually virtual registers, we may also use these two integers the source and destination registers in a new **##copy** instruction, which is then returned. On the other hand, **useful-instruction** saves the instruction's information in the expression graph by setting the appropriate values in `vregs>vns`, `exprs>vns`, and `vns>insns`. Note that in definitions using the syntax `:: word-name (stack -- effect) ... ;`, the input values in the stack effect are actually named lexical variables, like in most programming languages. Furthermore, the first line `insn dst>> :> vn` assigns the input instruction's destination register to the variable `vn`, which is used later in the definition of **useful-instruction**.

move this
to primer?

The **##copy** method of **process-instruction** cannot do anything to simplify the instruction, but will set the value number of the destination register to that of the source. By calling `vreg>vn` on the source register, we make sure to call **set-vn** between the destination and the canonical value number of the source.

Finally, the **array** method is used for the purposes of recursion, in the case that **rewrite** returns a sequence of replacement instructions. The correct action is, of course, to descend into this new sequence of instructions with **process-instruction**.

Underlying all of the redundancy elimination is the **rewrite** generic word. It has too many methods to look at the source code in-depth here, but it's instructive to give an overview of the transformations. These methods actually make up the bulk of the `compiler.cfg.value-numbering` code. They're spread across various sub-vocabularies. `compiler.cfg.value-numbering.rewrite` defines the generic itself, along with a handful of utilities. The method for the most general instruction class, `insn`, is defined to unconditionally return `f`, meaning no rewriting is performed by default. That way, we need only define **rewrite** methods for more specific instruction classes to get specialized behavior.

`compiler.cfg.value-numbering.alien` contains methods that simplify nodes related to Factor's FFI. Most involve fusing together the results of intermediate arithmetic. The instructions that access raw memory (namely **##load-memory**, **##load-memory-imm**, **##store-memory**, and **##store-memory-imm**) tend to have inputs to perform address arithmetic. Each has slots for a `base` register containing an address and a literal `offset` from it. But if `base` is defined by an **##add-imm** instruction, we can just update the `offset`, incrementing it by the literal operand of the **##add-imm**. Then, `base` will just be changed to the register operand of the **##add-imm**. This removes the memory instruction's need for the **##add-imm**, increasing the chances that the latter will become dead code to be removed later. Unlike the `-imm` variants, **##load-memory** and **##store-memory** also take a `displacement` register, which works like a non-immediate `offset`. Therefore, **##adds** can be similarly fused into **##load-memory-imm** and **##store-memory-imm** by transforming them into **##load-memory** and **##store-memory** instructions with the **##add**'s operand as the `displacement`. A few other similar transformations are also done, including rewrites for **##box-displaced-aliens** and **##unbox-any-c-pters**.

`compiler.cfg.value-numbering.comparisons` defines methods for the various branching and comparison instructions (which simply store booleans in registers, rather than branching upon them). The major optimizations performed are as follows:

- If possible, instructions are converted to more specific forms. For example, non-immediate instructions (e.g., `##compare`) may be turned into their `-imm` counterparts (e.g., `##compare-imm`) if one of their source registers corresponds to a literal value. `##compare-integer-imm` is also converted to `##test` if the architecture supports it. This corresponds to a special instruction in x86 that performs a bitwise AND for its side effects on particular flags, discarding the actual result. This can be more efficient when using the AND result as a boolean.
- If both inputs to a comparison or branch are literals, we may constant-fold the instruction. In the case of comparisons, this means converting it into a `##load-reference` of the proper boolean. In branches, this modifies the CFG so that the path which isn't taken is removed completely.
- Like a novice programmer writing `if (some_boolean != false) { ... }` in Java, the compiler may generate redundant boolean comparisons that need cleaning up. That is, the intermediate boolean values are eliminated when the result of a comparison is used by another comparison, collapsing the whole thing into a single instruction.

`compiler.cfg.value-numbering.folding` defines some auxiliary words for constant-folding arithmetic words. Mainly, `unary-constant-fold` and `binary-constant-fold` perform the actual operation on the one or two constant inputs provided. These words are used in `compiler.cfg.value-numbering.m`, which predictably simplifies math via standard rules. Arithmetic identities are rewritten—conceptually, $x + 0$ becomes just x , for instance. If self-inverting instructions (namely `##neg` for numerical negation and `##not` for boolean negation) are called on registers that themselves correspond to the same instruction, we can safely rewrite them into `##copy` instructions. Non-immediate instructions are converted to their `-imm` forms, if possible, and if both operands are constant, the expression is folded. The most interesting math optimizations use the associative and distributive laws. *Reassociation* conceptually converts $(x \otimes y) \otimes z$ into $x \otimes (y \otimes z)$ when both y and z are constants and \otimes is associative. So, for example,

```
##add-imm 1 X Y
##add-imm 2 1 Z
```

is converted into just

```
##add-imm 2 X (Y+Z)
```

where `X` is a virtual register, and `Y` and `Z` are constants. *Distribution* converts $(x \oplus y) \otimes z$ into $(x \otimes z) \oplus (y \otimes z)$, where y and z are constants, \oplus corresponds to addition or subtraction, and \otimes to multiplication or left bitwise shifts. Therefore,

```
##add-imm 1 X Y
##mul-imm 2 1 Z
```

is converted into

```
##mul-imm 3 X Y
##add-imm 2 3 (Y*Z)
```

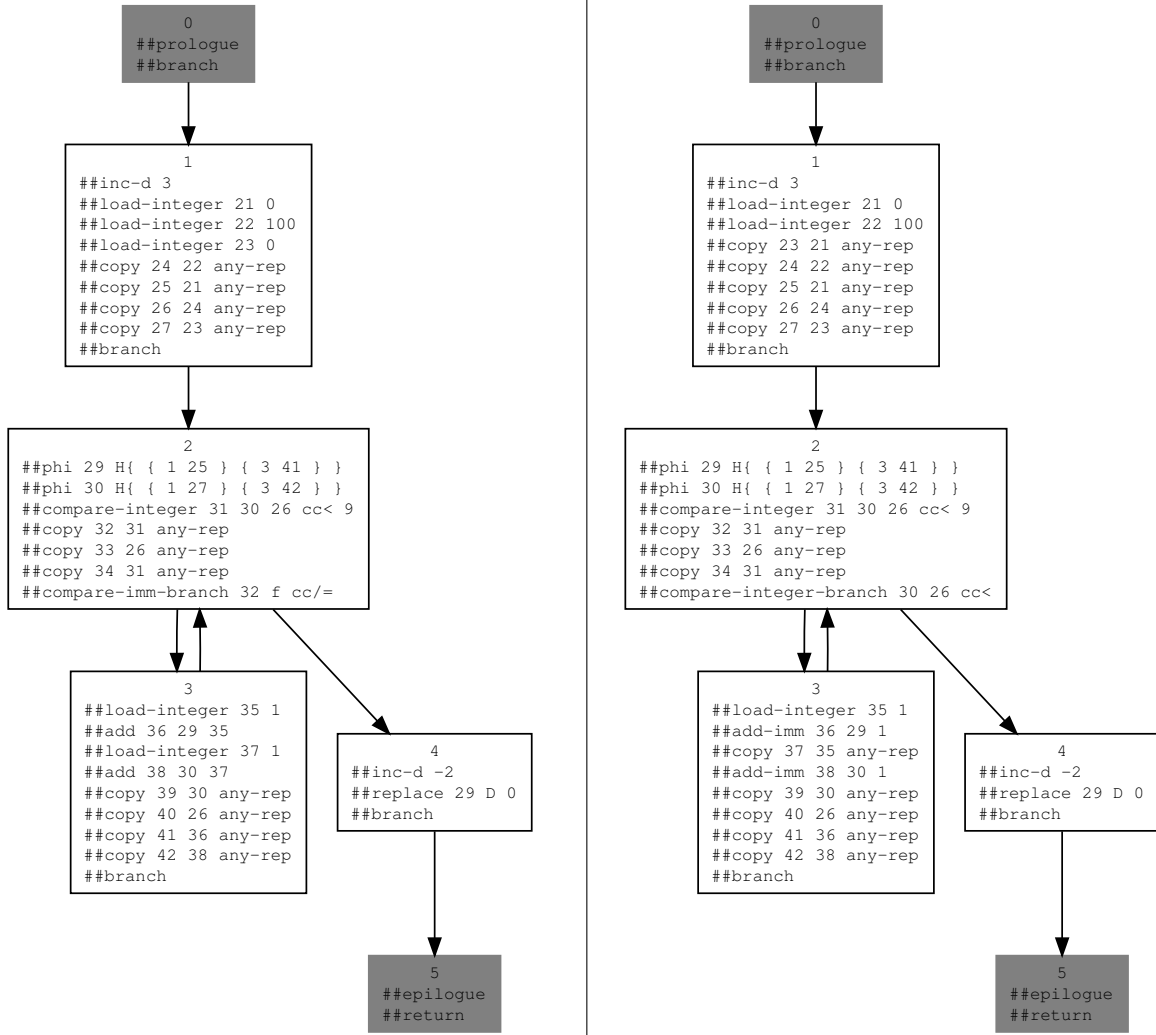


Figure 15: 0 100 [1 fixnum+fast] times before and after value-numbering

Notice that a new intermediate virtual register, 3, had to be created. However, if the product of Y and Z can be computed at compile-time and fits in an immediate operand, then we save cycles by using `##mul-imm` on a smaller number.

The last few methods of `rewrite` provide some obvious simplifications. `compiler.cfg.value-numbering.simd` performs some limited constant-folding for vector operations. `compiler.cfg.value-numbering.slots` propagates `##add-imm` address calculation to `##slot`, `##set-slot`, and `##write-barrier` instructions in a manner similar to `compiler.cfg.value-numbering.alien`. Finally, `compiler.cfg.value-numbering.m` provides a single method to rewrite `##replace` into `##replace-imm` if possible.

make sure margins don't get overflowed by long vocab names

make sure side-by-side figures aren't too crammed in final draft

To finish the discussion of local value numbering and Factor's particular implementation, we'll

examine the example from Figure 11 on page 37 in depth. For convenience, the before/after snapshot of the CFG is reproduced in Figure 15 on the previous page.

`value-numbering-step` begins at block 1, where `process-instruction` is **mapped** across the instructions. `##inc-d 3` does not have a `rewrite` method, so remains untouched; it is also not a `foldable-insn`, so it is simply returned. While `##load-integer 21 0` doesn't have a `rewrite` method, it is a `foldable-insn`, so `process-instruction` calls `check-redundancy`. At this point, the expression graph is empty. Calling `>expr` converts this instruction into an `integer-expr` object representing 0. `useful-instruction` leaves the tables as follows:

```
! vregs>vns
H{ { 21 21 } }

! exprs>vns
H{ { T{ integer-expr { value 0 } } 21 } }

! vns>insns
H{
  { 21 T{ ##load-integer { dst 21 } { val 0 } { insn# 1 } } }
}
```

The next instruction in block 1, `##load-integer 22 100`, behaves similarly, leaving:

```
! vregs>vns
H{ { 21 21 } { 22 22 } }

! exprs>vns
H{
  { T{ integer-expr { value 0 } } 21 }
  { T{ integer-expr { value 100 } } 22 }
}

! vns>insns
H{
  { 21 T{ ##load-integer { dst 21 } { val 0 } { insn# 1 } } }
  {
    22
    T{ ##load-integer { dst 22 } { val 100 } { insn# 2 } }
  }
}
```

The following instruction is `##load-integer 23 0`. In calling `check-redundancy`, we discover that the integer expression for 0 is already in `exprs>vns`, so this is turned into a `##copy`, and the value

number is noted. The remaining instructions in block 1 (aside from **##branch**) are all instances of **##copy**. **process-instruction** thus only sets their value numbers in the **vregs>vns** table, leaving them with the following at the end of block 1:

```
! vregs>vns
H{
    { 21 21 }
    { 22 22 }
    { 23 21 }
    { 24 22 }
    { 25 21 }
    { 26 22 }
    { 27 21 }
}

! exprs>vns
H{
    { T{ integer-expr { value 0 } } 21 }
    { T{ integer-expr { value 100 } } 22 }
}

! vns>insns
H{
    { 21 T{ ##load-integer { dst 21 } { val 0 } { insn# 1 } } }
    {
        22
        T{ ##load-integer { dst 22 } { val 100 } { insn# 2 } }
    }
}
```

Next, block 2 in Figure 15 on page 48 is processed. The tables are all reset, so even though block 1 happens to dominate block 2, none of its definitions are known to value-numbering. The **##phis** are ignored, as no important methods dispatch upon them. In trying to rewrite the **##compare-integer**, we call **vreg>vn** on the operands. Since they aren't in the **vregs>vns** table yet, they are assumed to be unique values. This assumption is pessimistic—we'd rather the values be the same, so we can remove redundancy. It happens to be correct here, though, as 26 corresponds to the integer 100, while 30 is an induction variable of the loop. However, **##compare-integer** cannot be rewritten into an immediate form, since our focus is local to the basic block, so we don't know that 26 has the value 100. The **##copy** instructions are processed as usual, and **##compare-imm-branch 32 f cc/=** is rewritten into a **##compare-integer-branch**, as the virtual register 32 has the same value (through the copies) as the **##compare-integer** result. This is a case of simplifying the **if (some_boolean != false) { ... }** pattern, and the definition of the register 31 becomes dead code after **rewrite** finishes with this last instruction. The expression graph is populated thus by the end:

yay, another term

```

! vregs>vns
H{
    { 32 31 }
    { 33 26 }
    { 34 31 }
    { 26 26 }
    { 30 30 }
    { 31 31 }
}

! exprs>vns
H{ { { ##compare-integer 30 26 cc< } 31 } }

! vns>insns
H{
    {
        31
        T{ ##compare-integer
            { dst 31 }
            { src1 30 }
            { src2 22 }
            { cc cc< }
            { temp 9 }
            { insn# 2 }
        }
    }
}

```

Once again, the tables are reset and we proceed to block 3. The first instruction, **##load-integer 35 1**, is entered into the expression graph. Since 35 is an operand of **##add 36 29 35**, **rewrite** changes this instruction into an **##add-imm**, as we know the constant value of the operand. The next **##load-integer** gets turned into a **##copy**, like in block 1, and the next **##add** is similarly changed to **##add-imm**. The copies do little but set more value numbers. As **process-instruction** calls **vreg>vn** on their sources, we'll insert entries into **vregs>vns** for those defined outside of the block, like 26. This leaves us with the following tables:

```

! vregs>vns
H{
    { 35 35 }
    { 36 36 }
    { 37 35 }
    { 38 38 }
    { 39 30 }
    { 40 26 }
}

```

```

    { 41 36 }
    { 26 26 }
    { 42 38 }
    { 29 29 }
    { 30 30 }
}

! exprs>uns
H{
    { { ##add-imm 30 1 } 38 }
    { { ##add-imm 29 1 } 36 }
    { T{ integer-expr { value 1 } } 35 }
}

! uns>insns
H{
    { 36 T{ ##add-imm { dst 36 } { src1 29 } { src2 1 } } }
    { 38 T{ ##add-imm { dst 38 } { src1 30 } { src2 1 } } }
    { 35 T{ ##load-integer { dst 35 } { val 1 } { insn# 0 } } }
}

```

The fourth invocation of `value-numbering-step` does not do anything interesting, as the `##replace` cannot be changed into a `##replace-imm`.

In summary, we managed to replace redundancies within basic blocks online by maintaining some simple hash tables. After copy propagation and dead code elimination, the CFG gets finalized to the one shown in Figure 16 on the next page. Because the value numbering algorithm was local, the `##compare-integer-branch` in block 2 could not be simplified to a `##compare-integer-imm-branch`, and we instead have to waste a register on the integer 100. But it's important to note that even considering a topological ordering of the CFG wouldn't have worked, as we'd have to ignore back-edges. The `##phis` that used to be in block 2 had inputs that flowed along the back-edge, and our pessimistic assumption would have to classify these values as distinct. One is for the counter introduced by `times`, and the other is from the top value of the stack being incremented by `fixnum+fast`. In this case, however, these induction variables are actually equal: both start at 0 and are incremented by 1 on each loop. In terms of the CFG in Figure 16 on the following page, the `EAX` and `EDX` registers are equivalent. Yet the combination of the pessimism and locality of the algorithm keep us from discovering this.

4.2 Global Value Numbering

Answering the challenges of Cocke, AWZ described what would be the de facto value numbering algorithm for several years, and rightly so. It was a properly *global* value numbering algorithm, working across an entire CFG instead of on single basic blocks. Their paper was important in another very relevant way: it is the first published reference to SSA form, including an algorithm for its construction.

cite-like

cite-like

cite Van-Drunen

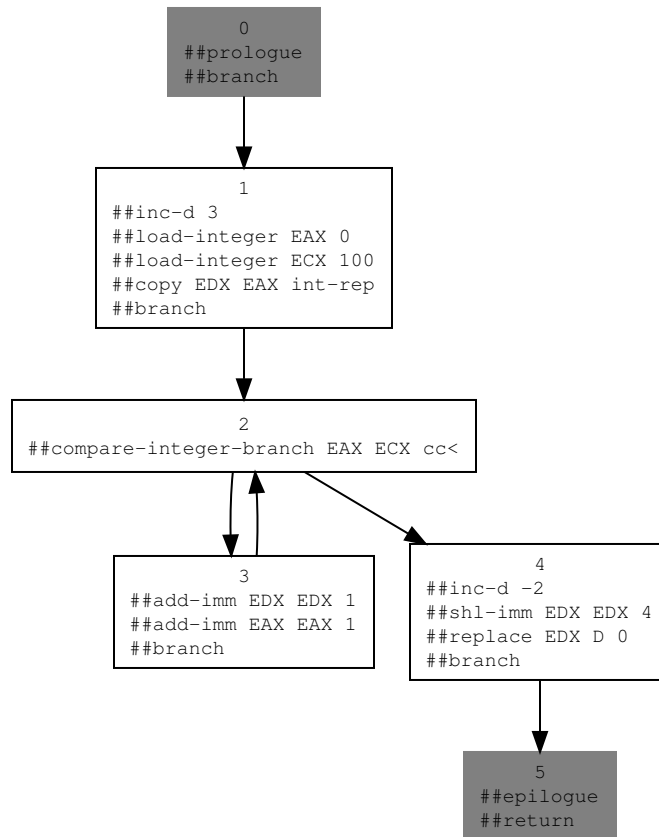


Figure 16: The final representation for 0 100 [1 fixnum+fast] times

Though we could try to extend the scope of Factor’s local value numbering, it is still inherently pessimistic. The algorithm of AWZ, which is commonly referred to simply as AWZ, uses a modification of Hopcroft’s minimization algorithm for finite state automata. It works on an *optimistic* assumption by first assuming every value has the same value number, then trying to prove that values are actually different. It does this by treating value numbers as *congruence classes* that partition the set of virtual registers. If two values are in the same class, then they are congruent, where congruence is defined as in Section 4.1.

cite-like

gls?

cite-like

Such a partition is not unique, in general. For instance, a trivial one places each value in its own congruence class. So, we look for the *maximal fixed point*—the solution that has the most congruent values and therefore the fewest congruence classes. We must start with a congruence class for each operation so that, say, all values computed by `##adds` are grouped together, those computed by `##muls` are in the same class, etc. We must then iteratively look at our collection of classes, separating them when we discover incongruent values. For an SSA variable in class P , we look at its defining expression. If an operand at position i belongs to class Q , then the i^{th} operand of every other value in P should also be in Q . Otherwise, P must be *split* by removing those variables whose i^{th} operands are not in class Q and placing them in a new congruence class. We keep splitting classes until the partitioning stabilizes.

The optimistic assumption may seem dangerous. Is it possible that we’re “overoptimistic”? That two values assumed to be congruent and not proven incongruent might actually be inequivalent when the program is run? The AWZ paper dedicates a section to proving that two congruent variables are equivalent at a point p in the program if their definitions dominate p . The proof is a bit quirky, but reasonable. They develop a dynamic notion of dominance in a running program which implies static dominance in the code, then show that congruence implies runtime equality (though equivalence does not imply congruence).

gls?

AWZ made the need for GVN algorithms apparent. However, finite state automata minimization makes for a more complicated algorithm than hash-based value numbering. A naïve implementation can be quadratic, although careful data structure and procedure design can make it $O(n \log n)$. Furthermore, it’s resistant to the same improvements we easily added to the local value numbering. To even consider the commutativity of operations requires changes in operand position tracking and splitting—the heart of the algorithm. It is generally limited by what the programmer writes down: deeper congruences due to, say, algebraic identities can’t be discovered.

gls?

In fact, by performing an optimization that uses the GVN information, more GVN congruences may arise. If we can somehow perform the two analyses simultaneously, they’ll produce better results. This generalizes to interdependent compiler optimizations at large, as elucidated in Click’s dissertation, which describes a method for formalizing and combining separate optimizations that make optimistic assumptions (whatever they happen to be for each particular analysis). He uses this to merge GVN with *conditional constant propagation*, which itself is a combination of constant propagation and unreachable code elimination (pretty much like the `propagate` pass from Section 3.2). Furthermore, GVN is extended to handle algebraic identities, propagate constants, and fold redundant ϕ s. Unfortunately, the straightforward algorithm for this is $O(n^2)$, while the $O(n \log n)$ version presented is not just complicated, but can also miss some congruences between ϕ -functions.

cite-like

cite Click,
Simpson

Hot on the heels of this work, Simpson’s dissertation provides probably the most exhaustive treatment of GVN algorithms. He presents several extensions, such as incorporating hash-based

cite-like

local value numbering into SSA construction, handling commutativity in AWZ GVN, and performing redundant store elimination. He builds off of the two classical algorithms independently, which underlines their inherent differences and limitations. In general, hash-based value numbering is easy to extend without greatly impacting the runtime complexity, as is the case in Factor's implementation.

gls?

Drawing from this experience, Simpson's hallmark algorithm combines the best of both worlds by taking the hash-based algorithm which is easy to understand, implement, and extend, and making it global, so it identifies more congruences. Dubbed the "reverse postorder (RPO) algorithm", it simply applies hash-based value numbering iteratively over the CFG until we reach the same fixed point computed by AWZ. (The fact that it computes the *same* fixed point is proven fairly straightforwardly in the dissertation.) It could technically traverse the CFG in any topological order, but Simpson defaults to reverse postorder.

gls?

Because it is based off the hashing algorithm, we get the benefits essentially for free. The same simplifications can be performed, but with the added knowledge of global congruences. Since the majority of Factor's value numbering code is dedicated to the `rewrite` generic, it makes sense to reuse as much of that code as possible. Therefore, to convert Factor's local algorithm to a global one, I modified the existing code to use the RPO algorithm.

The most fundamental change is to the expression graph. Referring to Listing 30 on the next page, we see most of the same code as in Listing 27 on page 43, with changes indicated by arrows (\longrightarrow). Two more global variables have been added, namely `changed?` and `final-iteration?`. The former is what we use to guide the fixed-point iteration. As long as value numbers are changing, we keep iterating. An important side effect of this is that we can no longer perform `rewrite` online, since the transformations we make aren't guaranteed to be sound on any iteration except the final one. This makes the RPO algorithm work *offline*, first discovering redundancies, then eliminating them in a separate pass. When this elimination pass starts, we'll set `final-iteration?` to `t`.

A key change is in the `vreg>vn` word, which now makes an optimistic assumption about previously unseen values. Given a new virtual register that wasn't in the `vregs>vns` table, the old version would map the register to itself, making the value its own canonical representative. However, if this version tries to look up a key that does not exist in the hash table, it will simply return `f` (which Factor will do by default with the `at` word). Therefore, every value in the CFG starts off with the same value "number", `f`. By the end of the GVN pass, there should be no value left that hasn't been put in the `vregs>vns` table, as we'll have processed every definition.

To keep track of whether `vregs>vns` changes, we simply need to alter `set-vn`. Here, we use `maybe-set-at`, a utility from the `assocs` vocabulary. This works like `set-at`, establishing a mapping in the hash table. In addition, it returns a boolean indicating change: if a new key has been added to the table, we return `t`. Otherwise, we return `f` only in the case where an old key is mapped to a new value. If an old key is mapped to the same value that's already in the table, `maybe-set-at` returns `f`. Therefore, when `vregs>vns` does change, we set `changed?` to `t` (which is what the `on` word does).

Finally, we define a new utility word, `clear-exprs`, which resets the `exprs>vns` and `vns>insns` tables. Unlike the local value numbering phase, we don't reset the entire expression graph. Instead, we make a pass over the whole CFG at a time. The only reason optimism works is that we keep trying to disprove our foolhardy assumptions. Really, `vregs>vns` establishes congruence classes of

```

→ ! Copyright (C) 2008, 2010 Slava Pestov, 2011 Alex Vondrak.
→ ! See http://factorcode.org/license.txt for BSD license.
USING: accessors kernel math namespaces assocs ;
→ IN: compiler.cfg.gvn.graph

SYMBOL: input-expr-counter

! assoc mapping vregs to value numbers
! this is the identity on canonical representatives
SYMBOL: vregs>vns

! assoc mapping expressions to value numbers
SYMBOL: exprs>vns

! assoc mapping value numbers to instructions
SYMBOL: vns>insns

→ ! boolean to track whether vregs>vns changes
→ SYMBOL: changed?

→ ! boolean to track when it's safe to alter the CFG in a rewrite
→ ! method (i.e., after vregs>vns stops changing)
→ SYMBOL: final-iteration?

: vn>insn ( vn -- insn ) vns>insns get at ;

→ : vreg>vn ( vreg -- vn ) vregs>vns get at ;

→ : set-vn ( vn vreg -- )
→     vregs>vns get maybe-set-at [ changed? on ] when ;

: vreg>insn ( vreg -- insn ) vreg>vn vn>insn ;

→ : clear-exprs ( -- )
→     exprs>vns get clear-assoc
→     vns>insns get clear-assoc ;

: init-value-graph ( -- )
    0 input-expr-counter set
    H{ } clone vregs>vns set
    H{ } clone exprs>vns set
    H{ } clone vns>insns set ;

```

Listing 30: The compiler.cfg.gvn.graph vocabulary

value numbers. At first, every value belongs in one class, **f**. We make a pass over the CFG to disprove whatever we can about this. If we've introduced new congruence classes (new values in the `vregs>vns` hash), we do another iteration. But each time, we use the congruence classes discovered from the previous iteration. At the start of each new pass, the expressions and instructions in `exprs>vns` and `vns>insns` are invalidated—their results are based on old information. So, these are erased on each iteration. Much like AWZ, we keep splitting classes until they can't be split anymore.

gls?

```
: value-numbering-step ( insns -- )
  [ simplify value-number ] each ;

: value-numbering-iteration ( cfg -- )
  clear-exprs [ value-numbering-step ] simple-analysis ;

: determine-value-numbers ( cfg -- )
  final-iteration? off
  init-value-graph
  '[
    changed? off
    _ value-numbering-iteration
    changed? get
  ] loop ;
```

Listing 31: Main logic in `compiler.cfg.gvn`

This logic is captured in Listing 31. Rather than reset the tables when we start processing each basic block in `value-numbering-step` like before, we call `clear-exprs` on each iteration over the CFG in `value-numbering-iteration`. Note that `value-numbering-step` no longer returns the changed instructions, as we aren't replacing them online. `value-numbering-iteration` uses `simple-analysis` instead of `simple-optimization`, which only expects global state to change—no instructions are updated in the block. Much to our advantage, `simple-analysis` already traverses the CFG in reverse postorder, so we needn't worry about traversal order. The top-level word `determine-value-numbers` ties this all together by calling `value-numbering-iteration` until we can get through it with `changed?` remaining false.

```
GENERIC: simplify ( insn -- insn' )

M: insn simplify dup rewrite [ simplify ] [ ] ?if ;
M: array simplify [ simplify ] map ;
M: ##copy simplify ;
```

Listing 32: Iterated rewriting in `compiler.cfg.gvn`

Note that the work of `value-numbering-step` is further divided into two words, `simplify` and `value-number`. These combine to do much the same work as `process-instruction` in List-

```

GENERIC: value-number ( insn -- )

M: array value-number [ value-number ] each ;

: record-defs ( insn -- ) defs-vregs [ dup set-vn ] each ;

M: alien-call-insn value-number record-defs ;
M: ##callback-inputs value-number record-defs ;

M: ##copy value-number [ src>> vreg>vn ] [ dst>> ] bi set-vn ;

: redundant-instruction ( insn vn -- )
  swap dst>> set-vn ;

:: useful-instruction ( insn expr -- )
  insn dst>> :> vn
  vn vn set-vn
  vn expr exprs>vns get set-at
  insn vn vns>insns get set-at ;

: check-redundancy ( insn -- )
  dup >expr dup exprs>vns get at
  [ redundant-instruction ] [ useful-instruction ] ?if ;

M: ##phi value-number
  dup inputs>> values [ vreg>vn ] map sift
  dup all-equal? [
    [ drop ] [ first redundant-instruction ] if-empty
  ] [ drop check-redundancy ] if ;

M: insn value-number
  dup defs-vregs length 1 = [ check-redundancy ] [ drop ] if ;

```

Listing 33: Assigning value numbers in compiler.cfg.gvn

ing 29. `simplify` makes the repeated calls to `rewrite` until the instruction cannot be simplified further. Its definition is in Listing 32 on page 57. We then pass the simplified instruction to `value-number`, which is defined in Listing 33 on the previous page. This also has a similar structure to `process-instruction`. The main difference is that instructions are no longer returned (again, they aren't altered in place). So, the `array` method uses `each` instead of `map` to recurse into the results of `rewrite`.

A subtle change is necessary with the `alien-call-insn` and `##callback-inputs` methods. Whereas `process-instruction` merely skipped over certain instructions that could not be rewritten, here we don't have that luxury. We need to be careful to `set-vn` every virtual register that gets defined by any instruction. While making a pessimistic assumption, it didn't matter if we did this: any unseen value would be presumed important by `vreg>vn`. However, with the optimistic assumption, `vreg>vn` will give the impression that unseen values are all the same by returning `f`. Therefore, we simply record the virtual registers defined in instructions that may define one or more of them. Specifically, `alien-call-insn` and `##callback-inputs` are classes that correspond to FFI instructions.

The `##copy` method uses `set-vn` the same way as before. `redundant-instruction`, `useful-instruction`, and `check-redundancy` are also largely the same. These have just been tweaked to not return instructions.

```
M: ##phi >expr
  inputs>> values [ vreg>vn ] map
  basic-block get number>> prefix
  \ ##phi prefix ;
```

Listing 34: ϕ expressions in `compiler.cfg.gvn.expressions`

The `##phi` method in Listing 33 on the preceding page represents a major change. Before, `##phis` were left uninterpreted. Congruences between induction variables that flowed along back-edges would not be identifiable. But now, by checking for redundant `##phis`, we may reduce them to copies. Each `##phi` object has an `inputs` slot, which is a hash table from basic block to the virtual register that flows from that block. Thus, there is one input for each predecessor. The `values` of the hash will be the virtual registers that might be selected for the `dst` value. We look up the value numbers of these, removing all instances of `f` with the `sift` word. If all of the inputs are congruent, we can call `redundant-instruction`, setting the value number of the `##phi`'s `dst` to the value number of its first input (without loss of generality). The `all-equal?` word will return `t` if the sequence is empty (as it's vacuously true), so we must make sure not to call `first` on the sequence, since this will be a runtime error. If the sequence is empty, we needn't note the redundancy, as the `##phi`'s `dst` will already have the optimistic value number `f` anyway. Otherwise, we call `check-redundancy`. The purpose of this is to identify `##phis` that are equal to each other. Even if its inputs are incongruent, a `##phi` might still represent a copy of another induction variable. So that `check-redundancy` works, we also define a `>expr` method in `compiler.cfg.gvn.expressions`, as seen in Listing 34. Here, the expression is an array consisting of the `##phi` class word, the current basic block's number, and the inputs' value numbers. We include the basic block number because only `##phis` within the same block can be considered

equivalent to each other.

The final method in Listing 33 on page 58 defines the default behavior for `value-number`, which calls `check-redundancy` on the simplified instruction if it defines a single virtual register. Note that we separate the `alien-call-insn` and `##callback-inputs` logic from this, since they happen to define a variable number of registers. If particular instances define only one register, we still don't want to call `check-redundancy` on them, since they don't have a `dst` slot. To avoid calling `dst>>` and triggering an error in `useful-instruction`, we needed separate methods for the FFI classes.

With these changes, we can globally identify value numbers, including equivalences that arise from simplifying instructions (even though no replacements are actually done yet). To illustrate this, consider again the example `0 100 [1 fixnum+fast] times`, reproduced in Figure 17 on the next page. As the expression graph changes frequently in this new algorithm, instead of showing the literal hash tables we'll use a shorthand notation. Virtual registers will be integers, and to avoid confusion value numbers will be written in brackets, like $\langle n \rangle$. Then, we'll show `vreg>vn` mappings with the notation $n \rightarrow \langle n \rangle$, where n is the register and $\langle n \rangle$ is the value number. If there is a corresponding expression in `exprs>vns`, it will be denoted after the mapping, like $n \rightarrow \langle n \rangle$ (*expression*). With the expressions, the instructions in `vns>insns` are a bit redundant for understanding the value numbering process, so they will be elided. Any mappings to `f` will be elided, as they're understood to be implicit when a key is absent.

Might make separate figures of each block, for easier reference

`determine-value-numbers` starts the first iteration, which of course starts at basic block 1. `##inc-d` is a no-op, but the first two `##load-integers` are established as useful instructions. `##load-integer 23 0` is recognized as redundant, since at this point we know that 21 has the value 0. The `##copy` instructions all pile on value number mappings, leaving us with the following:

```

21 → ⟨21⟩  (0)
22 → ⟨22⟩  (100)
23 → ⟨21⟩
24 → ⟨22⟩
25 → ⟨21⟩
26 → ⟨22⟩
27 → ⟨21⟩

```

At iteration 1, basic block 2, the first `##phi` has inputs 25 (from block 1) and 41 (from block 3). The former has the value number $\langle 21 \rangle$, while the latter is still at `f`. We treat this value number much like a \top element, unifying it with the other input to give us the assumption that 29 will be a copy of 25. Thus, it gets the same value number. A similar choice happens for the second `##phi`. The instruction `##compare-integer 31 30 26 cc< 9` is an interesting case. Due to our optimistic assumptions thus far, we believe 30 is carrying the value 0, and that 26 is set to 100. Thus, this instruction gets constant-folded by `simplify` into `##load-reference 31 t`. The CFG isn't changed, but the expression graph reflects this belief. Later, this assumption will be invalidated. The following copies are processed as usual, with the distinct difference here

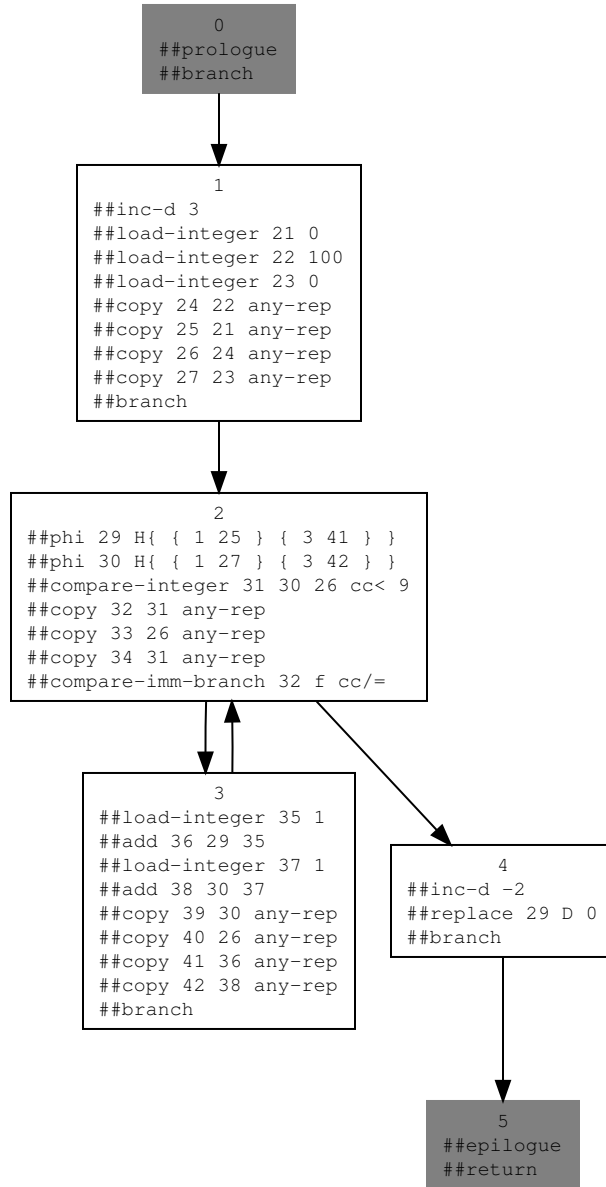


Figure 17: 0 100 [1 fixnum+fast] **times** before the new value numbering pass

that `##copy 33 26 any-rep` has the global knowledge of the value number of 26. Because the `##compare-integer` was constant-folded, so is the `##compare-imm-branch`—and to the same value, no less. This leaves us with:

```

21 → ⟨21⟩ (0)
22 → ⟨22⟩ (100)
23 → ⟨21⟩
24 → ⟨22⟩
25 → ⟨21⟩
26 → ⟨22⟩
27 → ⟨21⟩
29 → ⟨21⟩
30 → ⟨21⟩
31 → ⟨31⟩ (t)
32 → ⟨31⟩
33 → ⟨22⟩
34 → ⟨31⟩

```

Block 3 of iteration 1 gives the `##load-integers`' destinations the same value number, corresponding to the integer 1. Because optimism makes the algorithm think that 29 and 30 correspond to the integer 0, the `##adds` are constant-folded. This leaves us with:

```

21 → ⟨21⟩ (0)
22 → ⟨22⟩ (100)
23 → ⟨21⟩
24 → ⟨22⟩
25 → ⟨21⟩
26 → ⟨22⟩
27 → ⟨21⟩
29 → ⟨21⟩
30 → ⟨21⟩
31 → ⟨31⟩ (t)
32 → ⟨31⟩
33 → ⟨22⟩
34 → ⟨31⟩
35 → ⟨35⟩ (1)
36 → ⟨35⟩
37 → ⟨35⟩
38 → ⟨35⟩
39 → ⟨21⟩

```

40 \rightarrow $\langle 22 \rangle$
 41 \rightarrow $\langle 35 \rangle$
 42 \rightarrow $\langle 35 \rangle$

While block 4 is visited in each iteration, it doesn't define any registers, so doesn't affect the state of value numbering. Therefore, the above is the state left at the end of iteration 1.

Since `vregs>vns` clearly changed, iteration 2 commences by clearing the expressions, though the value numbers remain. Block 1 doesn't change from iteration 1, giving us:

21 \rightarrow $\langle 21 \rangle$ (0)
 22 \rightarrow $\langle 22 \rangle$ (100)
 23 \rightarrow $\langle 21 \rangle$
 24 \rightarrow $\langle 22 \rangle$
 25 \rightarrow $\langle 21 \rangle$
 26 \rightarrow $\langle 22 \rangle$
 27 \rightarrow $\langle 21 \rangle$
 29 \rightarrow $\langle 21 \rangle$
 30 \rightarrow $\langle 21 \rangle$
 31 \rightarrow $\langle 31 \rangle$
 32 \rightarrow $\langle 31 \rangle$
 33 \rightarrow $\langle 22 \rangle$
 34 \rightarrow $\langle 31 \rangle$
 35 \rightarrow $\langle 35 \rangle$
 36 \rightarrow $\langle 35 \rangle$
 37 \rightarrow $\langle 35 \rangle$
 38 \rightarrow $\langle 35 \rangle$
 39 \rightarrow $\langle 21 \rangle$
 40 \rightarrow $\langle 22 \rangle$
 41 \rightarrow $\langle 35 \rangle$
 42 \rightarrow $\langle 35 \rangle$

Now that we're in iteration 2, the inputs to the `##phis` of block 2 have been processed once before. For instance, we still believe that 25 corresponds to the integer 0 (which is incidentally correct), but now that 41 has the value number $\langle 35 \rangle$, we think it corresponds to the integer 1. While this is incorrect, it does break the congruence between the inputs, making the first `##phi` a useful instruction. The second `##phi`, however, still looks like a copy of the first. Even so, this is sufficiently different that the following `##compare-integer` cannot be constant-folded like before. However, it can still be converted to a `##compare-integer-imm`, as one of its operands corresponds to an integer. The redundant `##compare-imm-branch` gets rewritten to the same expression as the

`##compare-integer`, so winds up getting the same value number. This gives us:

```
21 → ⟨21⟩ (0)
22 → ⟨22⟩ (100)
23 → ⟨21⟩
24 → ⟨22⟩
25 → ⟨21⟩
26 → ⟨22⟩
27 → ⟨21⟩
29 → ⟨29⟩ (##phi 2 21 35)
30 → ⟨29⟩
31 → ⟨31⟩ (##compare-integer-imm 29 100 cc<)
32 → ⟨31⟩
33 → ⟨22⟩
34 → ⟨31⟩
35 → ⟨35⟩
36 → ⟨35⟩
37 → ⟨35⟩
38 → ⟨35⟩
39 → ⟨21⟩
40 → ⟨22⟩
41 → ⟨35⟩
42 → ⟨35⟩
```

Block 3 of iteration 2 also changes, since the `##adds` can't be constant-folded like before due to our new discovery about the `##phis`. However, the first one can still be converted to an `##add-imm`, and the second is marked the same as the first. This leaves the following value numbers:

```
21 → ⟨21⟩ (0)
22 → ⟨22⟩ (100)
23 → ⟨21⟩
24 → ⟨22⟩
25 → ⟨21⟩
26 → ⟨22⟩
27 → ⟨21⟩
29 → ⟨29⟩ (##phi 2 21 35)
30 → ⟨29⟩
31 → ⟨31⟩ (##compare-integer-imm 29 100 cc<)
32 → ⟨31⟩
```



```

33 → ⟨22⟩
34 → ⟨31⟩
35 → ⟨35⟩ (1)
36 → ⟨36⟩ (##add-imm 29 1)
37 → ⟨35⟩
38 → ⟨36⟩
39 → ⟨29⟩
40 → ⟨22⟩
41 → ⟨36⟩
42 → ⟨36⟩

```

Since the value numbers changed, we start iteration 3. The expressions are cleared, and block 1 once again does not change anything. The first **##phi** in block 2 still gets classified as useful, so no value numbers change. The major difference, though, is that the previous iteration's value numbers for registers in block 3 update the expression we have for the **##phi**. Whereas before we thought it was choosing between ⟨21⟩ (the integer 0) and ⟨35⟩ (the integer 1), the **##add** wasn't constant-folded in the previous iteration. Therefore, the virtual register 41 now corresponds to the result of the **##add** with the value number ⟨36⟩. We still can't disprove that the second **##phi** is different (because it, in fact, isn't). So, we're left with the following after iteration 3 finishes with block 2:

```

21 → ⟨21⟩ (0)
22 → ⟨22⟩ (100)
23 → ⟨21⟩
24 → ⟨22⟩
25 → ⟨21⟩
26 → ⟨22⟩
27 → ⟨21⟩
29 → ⟨29⟩ (##phi 2 21 36)
30 → ⟨29⟩
31 → ⟨31⟩ (##compare-integer-imm 29 100 cc<)
32 → ⟨31⟩
33 → ⟨22⟩
34 → ⟨31⟩
35 → ⟨35⟩
36 → ⟨36⟩
37 → ⟨35⟩
38 → ⟨36⟩
39 → ⟨29⟩
40 → ⟨22⟩

```

41 → ⟨36⟩
 42 → ⟨36⟩

Blocks 3 and 4 do not produce any more changes, so GVN has stabilized after 3 iterations, with our final congruence classes being:

⟨21⟩ = {21, 23, 25, 27}
 ⟨22⟩ = {22, 24, 26, 33, 40}
 ⟨29⟩ = {29, 30, 39}
 ⟨31⟩ = {31, 32, 34}
 ⟨35⟩ = {35, 37}
 ⟨36⟩ = {36, 38, 41, 42}

teletype the numbers in the align*s, I guess

4.3 Redundancy Elimination

Now that we’ve identified congruences across the entire CFG, we must eliminate any redundancies found. Since value numbering is now offline, this entails another pass. However, replacing instructions is more subtle with global value numbers than it is with local ones. Because values come from all over the CFG, we must consider if a definition is *available* at the point where we want to use it.

Figures 18 and 19 on pages 67–68 show the difference. In the former, we can see the CFG before value numbering for the code [10] [20] **if** 10 20 30. The two extra integers being pushed at the end are there to avoid branch splitting (see Section 3.3 on page 26). In block 4, there’s a **##load-integer** 27 10, which loads the value 10. In globally numbering values, we associate the **##load-integer** 22 10 in block 2 with the value 10 first, making it the canonical representative. However, we can’t replace the instruction in block 4 with **##copy** 27 22, because control flow doesn’t necessarily go through block 2, so the virtual register 22 might not even be defined. However, in Figure 19 on page 68, we see the CFG for the code 10 **swap** [10] [20] **if** 10 20 30. In this case, the first definition of the value 10 comes from block 1, which dominates block 4. So, the definition is available, and we can replace the **##load-integer** in block 4 with a **##copy**.

There are several ways to decide if we can use a definition at a certain point. For instance, we could use dominator information, so that if a definition in a basic block *B* can be used by any basic block dominated by *B*. However, here we’ll use a data flow analysis called *available expression analysis*, since it was readily implemented. Mercifully, Factor has a vocabulary that automatically defines data flow analyses with little more than a single line of code.

cite Simpson

Listing 35 on page 69 shows the vocabulary that defines the available expression analysis. It is a forward analysis based on the flow equations below:

cite?

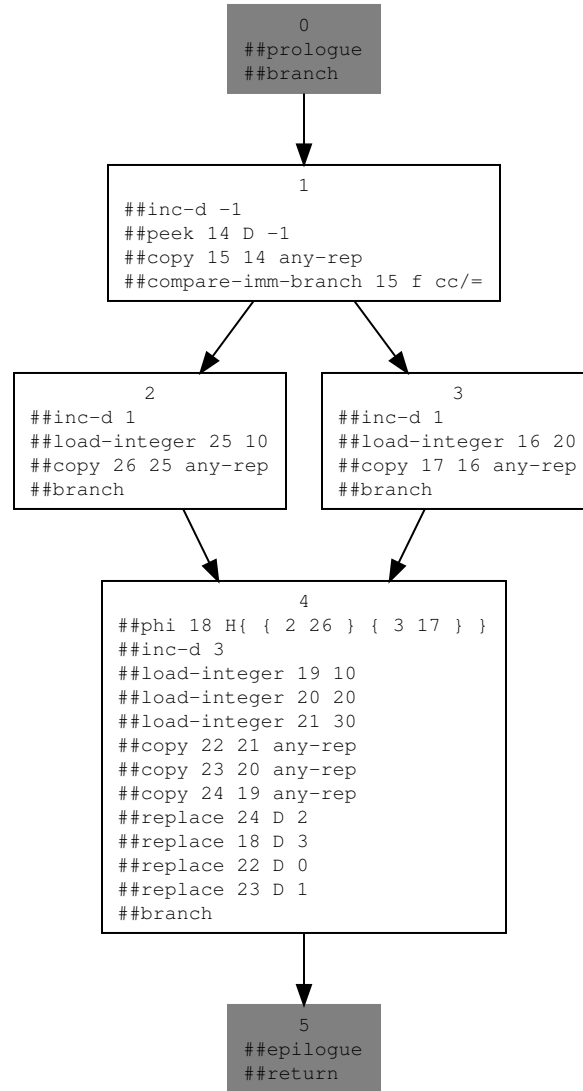


Figure 18: 10 is not available in block 4




Figure 19: 10 is available in block 4

```

! Copyright (C) 2011 Alex Vondrak.
! See http://factorcode.org/license.txt for BSD license.
USING: accessors assocs hashtables kernel namespaces sequences
sets
compiler.cfg
compiler.cfg.dataflow-analysis
compiler.cfg.def-use
compiler.cfg.gvn.graph
compiler.cfg.predecessors
compiler.cfg.rpo ;
FROM: namespaces => set ;
IN: compiler.cfg.gvn.avail

: defined ( bb -- vregs )
  instructions>> [ defs-vregs ] map concat unique ;

FORWARD-ANALYSIS: avail

M: avail-analysis transfer-set drop defined assoc-union ;

: available? ( vn -- ? )
  final-iteration? get [
    basic-block get avail-in key?
  ] [ drop t ] if ;

: available-uses? ( insn -- ? )
  uses-vregs [ available? ] all? ;

: with-available-uses? ( quot -- ? )
  keep swap [ available-uses? ] [ drop f ] if ; inline

: make-available ( vreg -- )
  basic-block get avail-ins get [ dupd clone ?set-at ] change-at ;

```

Listing 35: The compiler.cfg.gvn.avail vocabulary

$$\text{avail-in}_i = \begin{cases} \emptyset & \text{if } i = 0 \\ \bigcap_{j \in \text{pred}(i)} \text{avail-out}_j & \text{if } i > 0 \end{cases}$$

$$\text{avail-out}_i = \text{avail-in}_i \cup \text{defined}_i$$

Here, the subscripts indicate the basic block number. `definedi` denotes the result of the `defined` word from Listing 35 on the previous page. This returns the set of virtual registers defined in a basic block. Since we use virtual registers as value numbers, this is the same as giving us all the value numbers produced by a basic block. “Killed” definitions are impossible by the SSA property, so we needn’t track redefinitions of any virtual register. Using set intersection as the confluence operator means that the `avail-in` set will contain those values which are available on all paths from the start of the CFG to that block.

Using Factor’s `compiler.cfg.dataflow-analysis` vocabulary, the implementation takes all of two lines of code. The `FORWARD-ANALYSIS: avail` line automatically defines several objects, variables, words, and methods that don’t warrant full detail here. One we’re immediately concerned with is the `transfer-set` generic, which dispatches upon the particular type of analysis being performed and is invoked on the proper in-set and basic block. There is no default implementation, as it is the chief difference between analyses. So, the next line uses `defined` and `assoc-union` to calculate the result of the data flow equation. Other pieces we’ll see used are the top-level `compute-avail-sets` word that actually performs the analysis, the `avail-ins` hash table that maps basic blocks to their in-sets, and the `avail-in` word that is shorthand for looking up a basic block’s in-set.

We want to use the results of this analysis in the `rewrite` methods so that they won’t overstep their boundaries, and only make meaningful rewrites. However, we also want to use `rewrite` in the `determine-value-numbers` pass, where we don’t care about availability. In fact, we want to ignore availability altogether, so that we can discover as many congruences as possible. In order to separate these concerns, we need to have two modes for checking availability. Listing 35 on the preceding page defines the `available?` word to do just this. It will only check the actual availability if `final-iteration?` is true, otherwise defaulting to `t`. Therefore, during the value numbering phase, everything is considered available. We further define the utilities `available-uses?` and `with-available-uses?`. The former checks if all an instruction’s uses are available, and the latter does this only if another quotation first returns true. That way, we can guard instruction predicates with a test for available uses, like [`##add-imm?`] `with-available-uses?`.

Finding all the instances where `rewrite` needed to be altered was subtle. Since the old `value-numbering` was an online optimization, it didn’t need to worry about modifying an instruction in memory. But by doing the fixed-point iteration, we cannot permit `rewrite` to destructively modify any object instance until the final iteration. An obvious instance was in `compiler.cfg.value-numbering.comparisons` with the word `fold-branch`, responsible for modifying the CFG to remove an untaken branch. We definitely would not want the branch removed while doing the fixed-point iteration, because the transformation is not necessarily sound. So, we can protect it with a check for `final-iteration?`.

More typical were words like `self-inverse` from `compiler.cfg.value-numbering.math` (refer to Listing 37 on the next page). The idea is to change

```
##neg 1 2
##neg 3 1
```

```

! Before
: fold-branch ( ? -- insn )
  0 1 ?
  basic-block get [ nth 1vector ] change-successors drop
  \ ##branch new-insn ;

! After
: fold-branch ( ? -- insn )
  final-iteration? get [
    0 1 ?
    basic-block get [ nth 1vector ] change-successors drop
  ] [ drop ] if
  \ ##branch new-insn ;

```

Listing 36: Branch folding before and after

```

: self-inverse ( insn -- insn' )
  [ dst>> ] [ src>> vreg>insn src>> ] bi <copy> ;

! Before
M: ##neg rewrite
{
  { [ dup src>> vreg>insn ##neg? ] [ self-inverse ] }
  { [ dup unary-constant-fold? ] [ unary-constant-fold ] }
  [ drop f ]
} cond ;

! After
: self-inverse? ( insn quot -- ? )
  [ src>> vreg>insn ] dip with-available-uses? ; inline

M: ##neg rewrite
{
  { [ dup [ ##neg? ] self-inverse? ] [ self-inverse ] }
  { [ dup unary-constant-fold? ] [ unary-constant-fold ] }
  [ drop f ]
} cond ;

```

Listing 37: Rewriting words that are their own inverses

into

```
##neg 1 2
##copy 3 2 any-rep
```

since **##neg** undoes itself. But **rewrite** only has knowledge of one instruction at a time, so it looks up the redundant **##neg**'s source register in the **vregs>insns** table to see if it's computed by another **##neg** instruction. For straight-line code this is alright, but the source of the originating **##neg** (in the example, the virtual register 2) isn't necessarily available. So, we have to use **with-available-uses?** to make sure the virtual registers used by the result of a **vreg>insn** can themselves be used before we rewrite anything.

An even subtler issue that led to infinite loops occurred in simplifications like the arithmetic distribution in **compiler.cfg.value-numbering.math**. The problem is that the **rewrite** method would generate instructions that assigned to entirely brand new registers. These, of course, would invariably get value numbered, triggering a change in the **vregs>vns** table. A new iteration would begin, and (since it gets called on the same instructions as the previous iteration) **rewrite** would generate new virtual registers all over again. Therefore, the **vregs>vns** table would never stop changing. As a stop-gap, distribution had to be disabled altogether until the final iteration.

make it so

Armed with the correct rewrite rules, availability information, and global value numbers, we can perform global common subexpression elimination (GCSE). The logic in the **gcse** generic in Listing 38 on the following page is similar to **process-instruction** from Listing 29 on page 45 and **value-number** from Listing 33 on page 58. Unlike **value-number**, we do return an instruction (or sequence thereof) representing the replacement. Thus, the **array** method uses **map** instead of **each**, to hold onto the resulting sequence when recursing into several instructions.

defs-available is similar to **record-defs**, except that value numbers have already stabilized, so we don't call **set-vn**. Instead, we use the **make-available** word, which was the last one defined in Listing 35 on page 69. As we process the instructions of a block in order, we want to make sure that local rewrites can still be performed. We have to ensure that after processing an instruction, any register it defines is available to future instructions in the same block. Thus, we add the virtual register to that block's **avail-in** (which acts like a set, even though it's implemented by a hash table by Factor's data flow analysis framework). **alien-call-insns**, **##callback-inputs** instructions, and instances of **##copy** don't get rewritten any further, so we simply note that their definitions are available and move on.

The **?eliminate** word transforms an instruction into a **##copy** of the canonical value number that computes it. If the value number isn't available, we don't do anything but post-process with **defs-available**. If it is, a **##copy** is produced and its destination is made available. Thus, **eliminate-redundancy** works like **check-redundancy** from Listing 33 on page 58. We look up the expression computed by the instruction in the **exprs>vns** table. If it's there, we call **?eliminate**, but otherwise we leave the instruction alone and make its definitions available.

If the inputs to a **##phi** are all congruent, we'll call **?eliminate** to transform it into a **##copy** of its first input (without loss of generality). Otherwise, we check for equivalent **##phis** with **eliminate-redundancy**.

Finally, the **insn** method will default to calling **eliminate-redundancy** on instructions that define only one value, much like how **value-number** worked.


```

GENERIC: gcse ( insn -- insn' )

M: array gcse [ gcse ] map ;

: defs-available ( insn -- insn )
  dup defs-vregs [ make-available ] each ;

M: alien-call-insn gcse defs-available ;
M: ##callback-inputs gcse defs-available ;
M: ##copy gcse defs-available ;

: ?eliminate ( insn vn -- insn' )
  dup available? [
    [ dst>> dup make-available ] dip <copy>
  ] [ drop defs-available ] if ;

: eliminate-redundancy ( insn -- insn' )
  dup >expr exprs>vns get at
  [ ?eliminate ] [ defs-available ] if* ;

M: ##phi gcse
  dup inputs>> values [ vreg>vn ] map sift
  dup all-equal? [
    [ first ?eliminate ] unless-empty
  ] [ drop eliminate-redundancy ] if ;

M: insn gcse
  dup defs-vregs length 1 = [ eliminate-redundancy ] when ;

: gcse-step ( insns -- insns' )
  [ simplify gcse ] map flatten ;

: eliminate-common-subexpressions ( cfg -- )
  final-iteration? on
  dup compute-avail-sets
  [ gcse-step ] simple-optimization ;

```

Listing 38: Global common subexpression elimination in `compiler.cfg.gvn`

```

: value-numbering ( cfg -- cfg )
  needs-predecessors
  dup determine-value-numbers
  dup eliminate-common-subexpressions

  cfg-changed predecessors-changed ;

```

Listing 39: New global `value-numbering` word in `compiler.cfg.gvn`

The main loop works similarly to `determine-value-numbers`. `final-iteration?` is turned on (set to `t`), and we make sure to compute the `avail-in` sets needed to make `available?` work. Then, using `simple-optimization`, we iterate over each basic block. For each instruction, we first use `simplify` (refer to Listing 32 on page 57), then call `gcse` on the rewritten instruction. Thus, `rewrite` does the work of propagating available subexpressions and simplifying instructions, then `gcse` cleans up redundant instructions by converting them into `##copy` instructions if possible. The new `value-numbering` word can be seen in Listing 39.

oh no!
relative
time
stamps

Consider for the last time the example `0 100 [1 fixnum+fast] times`. Before, we had the CFG in Figure 20 on the next page. Making a final pass with `eliminate-common-subexpressions` gives us the CFG in Figure 21 on page 76. Compared to the CFG after the old `value-numbering` word was called (see Figure 11 on page 37), we have identified several more redundancies:

- The second `##phi` in block 2 has been turned into a `##copy` of the first.
- The `##compare-integer` of block 2 has been simplified to a `##compare-integer-imm`, since its operand 26 is both available and known to correspond to the integer value 100.
- Instead of a `##compare-integer-branch`, we similarly have a `##compare-integer-imm-branch` at the end of block 2.
- Because the `##phis` have been recognized as copies (i.e., the induction variables are congruent), the second `##add-imm` resulting from the old version of `value-numbering` is turned into a `##copy` of the first.

Afterwards, the `copy-propagation` pass cleans up all of these newly identified copies, as seen in Figure 22 on page 77. `eliminate-dead-code` now gets rid of more instructions than before, such as the second `##load-integer` in block 1, since it has been propagated to the `-imm` instructions in block 2. See Figure 23 on page 78. At last, after `finalize-cfg` in Figure 24 on page 79, we see a loop that uses a single register—down from the three in Figure 14 on page 40.

4.4 Results

The goal of improving the optimization in Factor is, of course, to reduce the average running time of programs, and to do so without changing their semantics. Short of formal verification, the latter requirement makes it necessary to thoroughly test any code that gets compiled with the new




Figure 20: 0 100 [1 fixnum+fast] **times** before the new value-numbering




Figure 21: 0 100 [1 fixnum+fast] **times** after the new value-numbering




Figure 22: 0 100 [1 fixnum+fast] **times** after copy-propagation




Figure 23: 0 100 [1 fixnum+fast] **times** after eliminate-dead-code




Figure 24: 0 100 [1 fixnum+fast] **times** after finalize-cfg

pass enabled. To this end, we'll employ Factor's extensive unit test coverage. Because Factor is (largely) self-hosting, its standard vocabularies are written in Factor code, typically coupled with tests. While some vocabularies will have more test coverage than others, the total amount of tests is quite large. By compiling each vocabulary and running their tests, we're indirectly testing the compiler: if tests that used to pass no longer do, then the new pass is changing the semantics of the code somehow. Though passing all tests does not guarantee the algorithm is correct, it does let us know that no known regressions have been introduced. Happily, with the new **value-numbering** phase enabled, all the same tests pass as before in a call to **test-all** from a freshly bootstrapped image.

The efficacy of the changes, on the other hand, must be measured relative to old benchmarks. Again, Factor has its bases covered, with a suite of 80 benchmarks run by the **benchmark** vocabulary. Each benchmark is run 5 times, where the garbage collector is run before each iteration. The minimum time from these runs is then used as the benchmark result. The data below comes from two separate runs of the **benchmarks** word, which invokes all the benchmark sub-vocabularies. The "before" time used the local value numbering, while "after" times had **value-numbering** replaced with the GVN pass. The "change" is measured by the formula

$$\frac{\text{before} - \text{after}}{\text{before}} \times 100$$

to indicate the relative running times. Negative values in this column are good, as that means the running time has decreased.

Guess I should provide my PC's specs

Benchmark	Before (seconds)	After (seconds)	Change (%)
benchmark.3d-matrix-scalar	3.705816738	3.046126696	-17.80
benchmark.3d-matrix-vector	0.161298778	0.089539887	-44.49
benchmark.backtrack	4.280001561	2.358672591	-44.89
benchmark.base64	5.127831493	2.853612485	-44.35
benchmark.beust1	7.531546384	4.604929188	-38.86
benchmark.beust2	20.308680548	12.843534349	-36.76
benchmark.binary-search	3.729776895	2.349520427	-37.01
benchmark.binary-trees	9.403166818	6.518867479	-30.67
benchmark.bootstrap1	32.472196349	30.887877896	-4.88
benchmark.chameneos-redux	2.923900422	2.041007328	-30.20
benchmark.continuations	0.273525202	0.200695972	-26.63
benchmark.crc32	0.010623653	0.005282642	-50.27
benchmark.dawes	1.588111926	1.027176578	-35.32
benchmark.dispatch1	7.640720326	5.106558985	-33.17
benchmark.dispatch2	5.221652668	3.984754032	-23.69
benchmark.dispatch3	9.710520454	6.203527737	-36.12
benchmark.dispatch4	8.224931156	4.098265543	-50.17
benchmark.dispatch5	4.74357434	3.478219608	-26.68
benchmark.e-decimals	3.903754723	2.646958072	-32.19
benchmark.e-ratios	4.774454589	3.658075473	-23.38

Benchmark	Before (seconds)	After (seconds)	Change (%)
benchmark.empty-loop-0	0.251816164	0.199189271	-20.90
benchmark.empty-loop-1	1.039242509	0.857588545	-17.48
benchmark.empty-loop-2	0.472215346	0.387974286	-17.84
benchmark.euler150	37.785852299	27.05450689	-28.40
benchmark.fannkuch	9.627490235	6.8970571	-28.36
benchmark.fasta	7.25292282	5.640517069	-22.23
benchmark.fib1	0.179389215	0.164933805	-8.06
benchmark.fib2	0.205853157	0.138174211	-32.88
benchmark.fib3	0.785036151	0.539739186	-31.25
benchmark.fib4	0.391805799	0.260370111	-33.55
benchmark.fib5	1.508625224	1.002724851	-33.53
benchmark.fib6	19.202504502	13.146010511	-31.54
benchmark.gc0	7.360087104	5.508594031	-25.16
benchmark.gc1	0.418173431	0.281497214	-32.68
benchmark.gc2	25.611210221	19.716168704	-23.02
benchmark.gc3	2.757943071	2.210785891	-19.84
benchmark.hashtables	8.068216942	7.997106348	-0.88
benchmark.heaps	4.360368411	4.32169158	-0.89
benchmark.iteration	7.875561986	6.277891729	-20.29
benchmark.javascript	17.881224721	12.74204052	-28.74
benchmark.knucleotide	5.490420772	3.5704101	-34.97
benchmark.mandel	0.251711276	0.198695557	-21.06
benchmark.matrix-exponential-scalar	16.451432774	12.017000042	-26.95
benchmark.matrix-exponential-simd	0.681684747	0.536850343	-21.25
benchmark.md5	10.40516678	9.198666403	-11.60
benchmark.mt	33.91981743	29.961085146	-11.67
benchmark.nbody	9.203478441	6.795154145	-26.17
benchmark.nbody-simd	0.845814208	0.854773096	+1.06
benchmark.nested-empty-loop-1	0.097090973	0.068475608	-29.47
benchmark.nested-empty-loop-2	0.893126911	0.861327078	-3.56
benchmark.nsieve	1.086110659	1.137648699	+4.75
benchmark.nsieve-bits	2.707271763	2.815509077	+4.00
benchmark.nsieve-bytes	0.785041878	1.211421146	+54.31
benchmark.partial-sums	3.762171661	4.130144177	+9.78
benchmark.pidigits	2.182877913	2.195385034	+0.57
benchmark.random	5.66540782	5.71913683	+0.95
benchmark.raytracer	5.047070171	4.39514879	-12.92
benchmark.raytracer-simd	1.072588515	0.980927338	-8.55
benchmark.recursive	2.703509403	2.529087637	-6.45
benchmark.regex-dna	2.208584014	1.808859571	-18.10
benchmark.reverse-complement	2.801163847	2.353254665	-15.99
benchmark.ring	1.822206473	1.62482491	-10.83
benchmark.sfmt	2.675838657	2.463367198	-7.94
benchmark.sha1	11.964973943	11.142380303	-6.88

Benchmark	Before (seconds)	After (seconds)	Change (%)
<code>benchmark.simd-1</code>	1.857778672	1.703206011	-8.32
<code>benchmark.sockets</code>	10.636346636	10.516448454	-1.13
<code>benchmark.sort</code>	0.695635429	0.581855635	-16.36
<code>benchmark.spectral-norm</code>	3.433630383	2.960833789	-13.77
<code>benchmark.spectral-norm-simd</code>	2.743240011	3.237017655	+18.00
<code>benchmark.stack</code>	1.580016742	2.004478602	+26.86
<code>benchmark.struct-arrays</code>	2.180774222	2.421915609	+11.06
<code>benchmark.sum-file</code>	0.883097981	0.957151577	+8.39
<code>benchmark.terrain-generation</code>	1.611800222	1.887047663	+17.08
<code>benchmark.tuple-arrays</code>	0.262747557	0.329399609	+25.37
<code>benchmark.typecheck1</code>	1.750223408	1.674592158	-4.32
<code>benchmark.typecheck2</code>	1.674738245	1.553203741	-7.26
<code>benchmark.typecheck3</code>	1.891206648	1.735390184	-8.24
<code>benchmark.ui-panes</code>	0.305595039	0.29985214	-1.88
<code>benchmark.xml</code>	3.013709363	2.722223892	-9.67
<code>benchmark.yuv-to-rgb</code>	0.398174487	0.318891664	-19.91

The results are promising: of 80 benchmarks, only 13 showed any increase in running time. And of those, even fewer showed significant increases. Duplicated below for convenience are the benchmarks that ran slower, sorted in decreasing order of the percent difference between running times. We can see the last five or six benchmarks exhibited negligible differences—not only is the relative change tiny, but the absolute difference in running times is less than 0.1 seconds. (The `benchmark.tuple-arrays` results also show a similar absolute change, but it is relatively much larger.)

Benchmark	Before (seconds)	After (seconds)	Change (%)
<code>benchmark.nsieve-bytes</code>	0.785041878	1.211421146	+54.31
<code>benchmark.stack</code>	1.580016742	2.004478602	+26.86
<code>benchmark.tuple-arrays</code>	0.262747557	0.329399609	+25.37
<code>benchmark.spectral-norm-simd</code>	2.743240011	3.237017655	+18.00
<code>benchmark.terrain-generation</code>	1.611800222	1.887047663	+17.08
<code>benchmark.struct-arrays</code>	2.180774222	2.421915609	+11.06
<code>benchmark.partial-sums</code>	3.762171661	4.130144177	+9.78
<code>benchmark.sum-file</code>	0.883097981	0.957151577	+8.39
<code>benchmark.nsieve</code>	1.086110659	1.137648699	+4.75
<code>benchmark.nsieve-bits</code>	2.707271763	2.815509077	+4.00
<code>benchmark.nbody-simd</code>	0.845814208	0.854773096	+1.06
<code>benchmark.random</code>	5.66540782	5.71913683	+0.95
<code>benchmark.pidigits</code>	2.182877913	2.195385034	+0.57

Overall, even transitioning to a relatively simple GVN algorithm amounts to a positive change in Factor’s compiler. More redundancies are eliminated, resulting in speedier programs. Judging

by unit tests, the implementation is at least as sound as the previous local value numbering, as all the same tests have passed.

4.5 Future Work

The GVN code presented in this thesis can be improved in various specific ways. Furthermore, the literature on GVN is extensive, and there are more overarching algorithmic strategies that have yet to be explored in the Factor code base. In this section, we explore some of these options for possible directions that Factor’s compiler can take from here.

As it stands, the new pass could be smarter. For instance, it does not consider the commutativity of certain operations. This would be straightforward to solve by making `>expr` sort the operands of commutative instructions’ expressions, thereby placing arguments in a canonical order. This would increase the number of congruences discovered between `##adds`, `##mults`, and even `##phis`. Also, the `copy-propagation` pass is remarkably similar to the new `value-numbering`—in fact, it could be removed altogether. All it does is collect global information about congruences as established by `##copy` instructions (by a similar fixed-point iteration), then replace the virtual registers of instructions with the original value (i.e., the one not established by a `##copy`). This allows `copy-propagation` to remove all `##copy` instructions. But the information calculated by `value-numbering` is a superset of this copy-equivalence data, so it should be easy to do global copy propagation in the GVN phase and save time on the redundant fixed-point iteration.

There remains an open question about the GVN implementation’s use of availability, too. As it stands, it’s rather strict: if the canonical value number for an expression is not directly available, `rewrite` gives up on reusing that value. However, the virtual registers which map to a single value number form a congruence class. We need not look just at the canonical leader (the first virtual register in the whole program to compute the particular expression). `rewrite` could change an instruction to reuse any member of the congruence class that was available. It remains to be seen when and if such a rewrite would be useful or desirable.

Existing literature also gives plenty of material for a better implementation. We can make the existing RPO algorithm more efficient in practice by observing that the only times we need to iterate are where there are cyclic dependencies between values in the CFG. For instance, the example from Section 4.3 only has cyclic dependencies in the induction variables: the `##phis` are defined by uses of virtual registers that are themselves defined by uses of the `##phi` targets in `##adds`. The RPO algorithm degenerates into the hash-based local algorithm of Section 4.1 on straight-line code. Thus, a more efficient algorithm will only iterate over the cycles between definitions instead of over the whole CFG.

Conceptually, we build a *value graph* (also known as *SSA graph*) whose nodes represent definitions and directed edges represent uses. Since it just codifies def-use information, we needn’t build an actual graph data structure. Using an algorithm due to Tarjan, SCCs of the value graph are iterated upon, while single nodes are processed only once. The strongly connected component (SCC) algorithm is more efficient than the RPO algorithm in practice, but the principles are the same. This gives us a comparatively simple, easily-extended GVN algorithm with complexity $O(nd)$, where n is the number of vertices in the value graph (i.e., the number of values we’re numbering) and d is the *loop connectedness* (the maximum number of back edges on any acyclic path)

cite

cite-like

of the value graph. Though d can theoretically be $O(n)$, in practice it seems to be bounded by a small constant. In Simpson's experiments, the maximum value of d was 4.

cite-like

A more thorough overhaul could incorporate further rewriting of the instructions. Gargi proposes a *predicated* value numbering algorithm that combines

cite-like

- optimistic value numbering, thus emulating Simpson's RPO and SCC algorithms;
- constant folding, algebraic simplification, and unreachable code elimination, thus emulating Click's strongest combination;
- global reassociation, thus performing the work already done in Factor;
- predicate inference, which can infer the values of comparisons dominated by some related predicate (i.e., comparisons in a block that is only reached via a particular conditional);
- value inference, which can infer the values of variables dominated by some related predicate (similar to range propagation, as seen in Section 3.2); and
- ϕ -predication, which (if possible) associates each input of a ϕ with the predicate that controls the path that leads to the selection of that argument, thus letting us find flow-sensitive congruences.

cite

cite

This combination is given in a *sparse* formulation, which makes it efficient enough to apply all of these optimizations. Essentially, when optimistic assumptions are invalidated (which, of course, happens as we iterate until reaching the fixed point), instead of recalculating every result (as in the RPO algorithm), we only recalculate the values that may yet change from this new information (as in the SCC algorithm).

Any portion of Gargi's algorithm may be selectively disabled, thus letting us tweak it for specific compile time vs. code quality trade-offs we might have. It promotes a fairly good separation of concerns in the algorithm, too, letting the pseudocode be presented piecemeal for each optimization. Gargi presents compelling examples of predicated value numbering's strength, and its addition to Factor could prove very worthwhile.

cite-like

cite-like

Historically, the development of GVN algorithms has forked along two concurrent bloodlines. Those we've studied thus far reflect the more "practical" line, which was largely implementation-driven and less formal than the other algorithms. But those focused on formal reasoning have recently become much more viable, and the wealth of ideas from them are worthwhile.

For those acquainted with Chapter 9 of The Dragon Book, this work will seem familiar, as it's rooted in the results of Kildall and Cousot, upon which the chapter is based. The former was a precursor to GVN, in that it described an algorithm for common subexpression elimination that partitioned expressions into congruence classes. However, its method was phrased in terms of *lattices*, which are algebraic structures that we can reason about formally. This is as in The Dragon Book: a lattice is a partially-ordered set for which any two elements have a unique *least upper bound* (or *join*) and *greatest lower bound* (or *meet*). By defining meet and join operators on a partially-ordered set of abstract values, we can represent many sorts of analyses on our programs.

cite

cite

cite

cite-like

Cousot formalize the salient properties of such interpretation over lattices in a framework

cite




Figure 25: Abstract interpretation over signs

dubbed *abstract interpretation*. To understand it intuitively, consider some arithmetic expression like (-5×14) . Our first inclination is probably to interpret it with respect to numeric values, but we can understand it in several different contexts. Let’s use signs (+, −, and ±) as our abstract domain and consider the operators to be defined by the rules of signs. Figure 25 shows a lattice we can use for this. Using a version of \times cast in the context of this lattice, we can interpret (-5×14) as

$$-5 \times 14 \rightarrow (-) \times (+) \rightarrow (-)$$

proving that (-5×14) is negative. Using this framework, the results are correct, but only useful within the confines of what we define. For instance, we can interpret $(-5 + 14)$ as

$$-5 + 14 \rightarrow (-) + (+) \rightarrow (\pm)$$

proving very little—the result is either positive or negative.

Despite the inherent limitations, we find the results useful as approximations of more complex properties. For example, we used congruence to approximate runtime equivalence. Only a year before AWZ was published, Steffen showed that Kildall’s approach could be framed as abstract interpretation over *Herbrand equivalences*—that is, equivalences where operators are uninterpreted. This is actually the same notion of congruence we had from before: expressions are equivalent if their operators and operands are equivalent, irrespective of the result of applying the operator.

gls?

cite

cite-like

The primary strength of the abstract interpretation approaches are that they are *complete*; intuitively, there is no loss of information at each step of abstract interpretation. However, this “loss of information” is relative to the information encompassed by the abstract domains. While we can find all Herbrand equivalences, we aren’t guaranteed to find equivalences induced by interpreting operators, which was effectively the work done by combining optimizations (e.g., constant folding is the interpretation of certain operators upon constant operands). So, while complete, these approaches vary in *preciseness*. Most of the work in the abstract interpretation of GVN did little to study the results of interpreting the same operators we saw before, but note it’s a promising direction for future research.

cite

The cost of this completeness has traditionally been exponential time complexities. There have been several attempts to remedy this. RKS note AWZ is incomplete, since it treats ϕ functions as

cite

gls?

uninterpreted, so fails to discover congruences between ϕ s and ordinary expressions. Their attempt to improve upon it alternately applies AWZ and the normalization rules

$$\begin{aligned}\phi(a \otimes b, c \otimes d) &\rightarrow \phi(a, c) \otimes \phi(b, d) \\ \phi(x, x) &\rightarrow x\end{aligned}$$

gls?

until the partitioning reaches a fixed point. However, this is $O(n^2 \log n)$ in the expected case— $O(n^4 \log n)$ in the worst case—and it turned out to be incomplete not just in the presence of cycles but also in certain acyclic code.

cite RKS

Later, Gulwani04 furthered the quest for an efficient, complete GVN algorithm in a novel way by using *randomized interpretation* (which is what it sounds like). The paper even explored various interpretations—specifically of linear combinations, bitwise operators, memory loads/stores, and integer division—that could make results more precise. But it was still $O(n^4)$ and ran a small chance of making incorrect inferences due to its randomized nature. For compilers, this isn't really acceptable, though such a scheme could be used in things like program verification tools.

cite Gulwani07

cite

cite Nie

cite

From their trip back to the drawing board, Gulwani07 returned with a polynomial time algorithm for GVN that is complete for all Herbrand equivalences among terms of a limited size. Choosing a size bound equal to the size of the entire program is clearly sufficient. Note, however, that this is specifically for Herbrand equivalences; they do not show their results for any interpreted operators, but note it's an important area for exploration. Adding to this, Nie present a similar algorithm, except based on SSA form. Both wind up using the same size restrictions to guarantee complexity. Both also use an additional special-purpose data structure to represent the set of Herbrand equivalences and to perform abstract evaluations over them, which adds a conceptual load to the algorithms and might make them more difficult to implement. However, unlike most other abstract interpretation-based algorithms, Nie's is demonstrably practical, as the authors implemented it for GCC. In their experiments, the size restriction turned out to be unnecessary for avoiding the exponential case, showing that the main bottleneck in complete GVN algorithms is typically their poor data structure choices.

cite

cite

gls?

Clearly, there is much room for future exploration. Even without crossing the boundaries of GVN into the scope of other compiler optimizations, we can eliminate all sorts of redundancies. The literature has a wealth of algorithms that all warrant experimentation. With varying degrees of aggressiveness, there are several opportunities to make Factor a more efficient high-level language.