

Figure 1: Visualizing stack-based calculation

1 Language Primer

citations for this history are fragmented across the internet; should consolidate some kernel of citation from it

Factor is a rather young language created by Slava Pestov in September of 2003. Its first incarnation targeted the Java Virtual Machine (JVM) as an embedded scripting language for a game. As such, its feature set was minimal. Factor has since evolved into a general-purpose programming language, gaining new features and redesigning old ones as necessary for larger programs. Today's implementation sports an extensive standard library and has moved away from the JVM in favor of native code generation. In this section, we cover the basic syntax and semantics of Factor for those unfamiliar with the language. This should be just enough to understand the later material in this thesis. More thorough documentation can be found via Factor's website.

figure out section / chapter names

cite chapter?

cite factor-code.org

1.1 Stack-Based Languages

Like Reverse Polish Notation (RPN) calculators, Factor's evaluation model uses a global *stack* upon which operands are pushed before operators are called. This lends itself to *postfix* notation, in which operators are written after their operands. For example, instead of $1 + 2$, we write $1\ 2\ +$. Figure 1 shows how $1\ 2\ +$ works conceptually:

- 1 is pushed onto the stack
- 2 is pushed onto the stack
- $+$ is called, so two values are popped from the stack, added, and the result (3) is pushed back onto the stack

Other stack-based programming languages include Forth, Joy, Cat, and PostScript.

cite

The strength of this model is its simplicity. Parsing becomes very flexible, since whitespace is basically the only thing that separates tokens. In the Forth tradition, functions (being single tokens delineated by whitespace) are called *words*. This also lends to the term *vocabulary* instead of “module” or “library”. In Factor, the parser works as follows.

cite

cite

cite

cite

- If the current character is a double-quote (`"`), try to parse ahead for a string literal.
- Otherwise, scan ahead for a single token.
 - If the token is the name of a *parsing word*, that word is invoked with the parser's current information.




Figure 2: Quotations

- If the token is the name of an ordinary (i.e., non-parsing) word, that word is added to the parse tree.
- Otherwise, try to parse the token as a numeric literal.

Words correspond to *first-class* functions from other languages. That is, words are themselves objects that just perform instructions whenever they’re told—most commonly, when Factor is executing a piece of code. While infix languages tend to distinguish syntactically significant operators (e.g., for arithmetic) from normal user-defined functions, there is no such difference here. Since `1 2 +` and `1 2 foo` are tokenized the same way, other languages’ special operators are just words in Factor. This simple syntax gives rise to a simple evaluation model that essentially goes left-to-right (as we saw in Figure 1 on the previous page): literals are pushed to the stack, words are executed.

Factor syntax examples; figure out listings (grrr)

Parsing words serve as hooks into the parser, letting Factor users extend its syntax dynamically. For instance, instead of having special knowledge of comments built into the parser, the parsing word `!` scans forward for a newline and discards any characters read (adding nothing to the parse tree).

Other parsing words act as sided delimiters. The parsing word for the left-delimiter will scan ahead for the right-delimiter, using whatever was read in between to add objects to the parse tree. For example, array literals are created by the parsing word `{` parsing ahead until it finds a `}` token, collecting the results into an array that’s added to the parse tree. Since it parses recursively, array literals can be nested, contain comments, etc. That is, anything parsing words do will work inside of an array. If an ordinary word is parsed, it is added to the array—not executed. Thus, `{ 1 2 + }` denotes an array of three elements: the numbers 1 and 2, and the word `+`.

example

offshore to
example
listing?

Figure out how to typeset stacks in a more consistent way

A particularly important set of parsing words in Factor are the square brackets, `[` and `]`. Any code in between such brackets is collected up into a single sequence called a *quotation*. Essentially, it’s a snippet of code whose execution is suppressed. The code inside a quotation can then be run with the `call` word. Quotations are like anonymous functions in other languages, but the stack model makes them conceptually simpler, since we don’t have to worry about variable binding and the like. Consider a small example like `1 2 [+] call`. You can think of `call` working by “erasing” the brackets around a quotation, so this example behaves just like `1 2 +`. Figure 2 shows its evaluation: instead of adding the numbers immediately, `+` is placed in a quotation, which is pushed to the stack. The quotation is then invoked by `call`, so `+` pops and adds the two numbers and pushes the result onto the stack. We’ll show how quotations are used in `\ref{sec:primer:control-flow}`.

ref

1.2 Concatenative Programming

The biggest step needed to understand Factor isn't really big: everything else follows from the stack-based structure outlined in Section 1.1. With some input on the stack, consecutive words transform the data in discrete steps, thereby shaping a result. This breeds a fairly stylized way of writing code that resembles a pipeline. Postfix code like `1 2 +` often seems to read “backwards”, but when viewed in this pipelined fashion, it can read more clearly than standard syntax. For instance, suppose we're given a value `x` to manipulate by applying three functions, `f`, `g`, and `h`, in order. Typical function application notation looks like `h(g(f(x)))`, which must be read from the inside out because of the nesting. In Factor, we'd write `x f g h`, which is flat and reads left-to-right. Having such a whitespace-centered syntax reduces noise: instead of grouping arguments, the order of evaluation is the order you write your code.

In a way, this makes whitespace equivalent to function composition: just treat every word as a function from stacks to stacks—“before” to “after”. Even literals (numbers, strings, etc.) can be thought of as functions that take in a stack and return that stack with an extra element pushed onto it. Factor is thus said to be a *concatenative* language: concatenating any two programs denotes the composition of the functions the programs individually denote. Note that being stack-based is unnecessary for being concatenative, though it's common.

cite and
talk about
Joy

list other
concatenative
languages

cite Joy

This functional interpretation is useful for formally reasoning about concatenative languages.

- Concatenative programming
 - Point-free style
 - Origin of name “Factor”
- Stack effects
 - Basic stack effects: stack shufflers illustrated
 - Complex stack effects: row polymorphism & types
 - Stack checker
- Control flow
 - Combinators
 - * if
 - * each
 - * while
 - Dataflow combinators
 - * Dip/keep
 - * Cleave
 - * Spread
 - * Apply
- Object system

- tuples
 - generics & methods
- Libraries & metaprogramming
 - Results of evolution
 - locals?
 - fry?
 - macros?
 - functors?
 - ffi?